

NON-INTERACTIVE SCHNORR THRESHOLD IDENTIFICATION

AUTHOR NAME

ABSTRACT. TODO native threshold protocol as much as the ordinary protocol
TODO We propose a novel approach to threshold identification based on the
Schnorr protocol. We provide evidence that the protocol has good chances
to be secure by demonstrating a formal security proof in the random oracle
model under one-more discrete-logarithm (OMDL) hardness.

1. INTRODUCTION

Distributing the prover of a zero-knowledge protocol is usually achieved by means of a trusted combiner. The shareholders engage in some distributed computation in order to generate a uniform commitment and then use it to generate ordinary proofs for their local shares. After aggregating the respective proof packets, the combiner couples them into an ordinary proof and forwards it to the verifier, who proceeds to verification like in the non-distributed case; in particular, the combined proof resides in the space of proofs that could have been generated for the combined secret directly. The idea is similar to that of the Schnorr threshold signature scheme [1]. It has been elaborated in the proof-of-knowledge context by [2], who apply it in order to distribute protocols as general as proving knowledge of the preimage of a homomorphism (cf. [8]).

The integration of such protocols is seamless on the verifier's side because it does not need to either trust or maintain the public keys of the shareholders. This convenience comes at significant cost on the prover's side. First and foremost, the distributed protocol is highly interactive even though its non-distributed counterpart may consist of a single communication step. This is because of the multi-party computation for the uniform commitment, usually a distributed key generation (DKG) scheme, which introduces significant network overhead upon every protocol round while also relying on real-time communication and subtle security assumptions regarding the involved parties. Observe that delegating combination to the verifier's side makes no difference (apart from fewer trust assumptions) since the protocol would remain equally interactive. Techniques for reducing interactivity, e.g., pseudorandom secret sharing, do not scale beyond a limited number of involved provers and, most importantly, require coordination or some relevant infrastructure to be maintained and constantly available. While this may not be an issue for monitored operations with long-term impact in a regulated environment (e.g., distributed signatures stored in a public ledger), it becomes prohibitive if the provers are compelled to act unattendedly and asynchronously in the absence of infrastructure, e.g., for the purpose of multi-factor authentication or when performing some joint IoT operation in order to rapidly release a sensitive action. We are primarily interested for situations of the latter kind.

A relevant paradigm is that of threshold decryption, which employs decoupled zero-knowledge proofs in order to detect maliciously fabricated partial decryptors. Given a ciphertext, the shareholders attach a Chaum-Pedersen proof to their respective decryptor, proving that this is securely bound to their secret share; after aggregating sufficiently many of these packets, the decrypting party reconstructs the combined decryptor only if the attached proofs verify against the respective public shares, i.e., decryption is not released unless all attached proofs are individually found to be valid. Since every Chaum-Pedersen proof includes a Schnorr proof for the respective shareholder’s key, one can easily eliminate the encryption specific details in order to derive a threshold identification protocol.

A protocol of the above kind is non-interactive and requires no coordination or infrastructure on the provers’ side, allowing the shareholders to act independently; however, it does require from the decrypting party (i.e., the verifier) to maintain the correct public shares in order to be able to individually verify the attached proofs. This evidently entails the need for a trusted multi-key setup in order to certify both the ownership of the public shares *and* their compatibility (i.e., that they indeed comprise a sharing of the group public key), otherwise the protocol remains susceptible to all sorts of rogue-key attacks. Observe that rogue-key attacks may not be a major concern for threshold decryption because proofs are verified against both the respective shares and the given ciphertext; if the ciphertext has been generated by the encrypting party with respect to the correct group public key, a rogue key will most probably be detected. This is not so for the derived identification protocol, where the verification operation involves nothing beyond the registered public shares. In this case, security becomes entirely dependent on the multi-key setup and its generally intractable attack surface, which can only be analysed per use case. Even if a sufficiently encompassing multi-key setup could be designed and analysed with respect to a specific attack model, this does not guarantee its adoptability in the real world; any such setup is not straightforward to integrate with and, if adopted, it comes at the cost of restricted interoperability and significant infrastructure extensions and maintenance. Evidently, it is completely useless in a flexible and unregulated identification setting, where no trust infrastructure can be assumed to be in place at all.

1.1. Contribution. The above discussion shows that non-reliance on the public shares and non-interactiveness on the provers’ side appear to be diametrically opposed properties; eliminating infrastructure on one side requires its existence and maintenance on the other. We define an infrastructure-free Schnorr-based threshold identification protocol which achieves non-interactiveness and non-reliance on the public shares at once; as such, it becomes applicable in flexible and unregulated decentralized settings and can be cheaply adopted in the following sense: (a) existing verifiers need only extend their software in order to support the distributed case, without adding infrastructure for trusting, registering and maintaining the shareholders’ public keys (b) existing keys can be shared at the verifier’s complete negligence. No trusted key setup is needed beyond the minimum assumption that some key has been shared amongst a group of shareholders according to some Shamir-based (n, t) -sharing (no matter how) and the group public key has been advertised in some “trusted” fashion; in particular, the protocol should be secure against impersonation attacks in the plain-public key model (cf. [12]). We provide evidence that this has good chances to hold true by giving a formal security proof in

the random oracle model under the one-more discrete-logarithm (OMDL) hardness assumption.

1.2. Organization. In Section 2 we get a closer look at distributed Schnorr identification and argue in detail why existing solutions fail to meet its needs in the real life, specifically with respect to unregulated and decentralized settings. In Section 3 we introduce the proposed threshold identification scheme, derive useful validity conditions and make some high-level security considerations, including a key-recovery attack against its interactive counterpart (cf. Section 3.2.3). In Section 4 we outline the general threat model regarding the security against impersonation attacks and define the relevant attack games. We do not care to formulate auxiliary notions for proof-of-knowledge schemes (e.g., soundness, honest-verifier zero-knowledge, witness indistinguishability etc.) in the threshold setting, focusing instead directly on security against impersonation attacks; however, we do specify a restricted notion of extractability (in the sense of special soundness) in order to derive a reasonable attack model under the OMDL hardness assumption. In Section 5 we present the formal security proof in the random oracle model; this is done by fork and extraction, using the key-recovery attack against the interactive protocol and the Local Forking Lemma. The formalism and main facts regarding Shamir’s secret sharing are presented in Appendix A; a quick refresher on the ordinary Schnorr protocol can be found below.

1.3. Schnorr protocol. Fix a cyclic group $\mathbb{G} = \langle g \rangle$ of order q . By secret key is meant a scalar $x \in \mathbb{Z}_q$, with the group element $y = g^x$ being its public counterpart. Equivalently, $x = \log y$, where the logarithm is taken with respect to the generator g of the underlying group structure. *Discrete-logarithm (DL) hardness* for the system under consideration is the assumption that computing the logarithm of a uniformly random element is infeasible for every existing computer. More accurately, given $y \leftarrow_{\$} \mathbb{G}$, every probabilistic polynomial-time algorithm should correctly compute $x \leftarrow \log y$ with negligible chance of success, where negligibility is determined by the current level of disposable computational power. In more applied terms, it should be impossible to infer a truly random key from its public counterpart within reasonable time. Industrial public-key cryptography relies increasingly more on the plausibility of this assumption over carefully chosen groups, a trend that will only be challenged with the advent of practical quantum computing.

The *Schnorr identification protocol* is a non-interactive zero-knowledge (NIZK) scheme, allowing the key holder of $y = g^x$ to prove knowledge of x without revealing any information about the latter. Its interactive variance is a sigma protocol. The prover samples $r \leftarrow_{\$} \mathbb{Z}_q$, stores it for the current session and sends the commitment $u = g^r$ to the verifier. The verifier samples a challenge $c \leftarrow_{\$} \mathbb{Z}_q$, caches it for the current session and sends it to the prover. Finally, the prover responds with $s = r + cx$ and the verifier accepts only if the relation

$$g^s = y^c u$$

is satisfied; indeed, $g^s = g^{r+cx} = (g^x)^c g^r = y^c u$. In the presence of a secure hash function $H : \mathbb{G} \rightarrow \mathbb{Z}_q$, the protocol can be made non-interactive (1-step) by means of the Fiat-Shamir transform. Specifically, the prover precomputes $c \leftarrow H(u)$ and sends (u, s) to the verifier, who retrieves c from u and proceeds to verification. Non-interactiveness is most often preferable due to greater robustness, reduced

throughput and non-reliance on real-time communication. The intuition behind is that a secure hash function should inject sufficiently high entropy across every bit of its output, so that assigning $c \leftarrow H(u)$ becomes indistinguishable from sampling c uniformly at random. The current mainstream formalizes this fact within the *random oracle model* (ROM).

The Schnorr protocol is an “identification” scheme because it verifies the holder of a key against its public identity in a sound and complete fashion. Note that key semantics are not essential; the protocol identifies any procedure that “sees” the secret logarithm of some given group element. It is namely a zero-knowledge (ZK) proof-of-knowledge primitive for the discrete logarithm and as such applicable in any context where this knowledge arises as crucial (e.g., in plain ElGamal encryption, where a proof for the involved randomness may be attached in order to derive a non-malleable ciphertext); most notably, it is the building block for proving knowledge of witnesses to arbitrary linear relations, common instances of which are the Chaum-Pedersen protocol for Diffie-Hellman tuples and the Okamoto protocol for representations. From this perspective, security against impersonation attacks amounts to the infeasibility of fabricating a proof that verifies against a group element without knowing its logarithm; in the proof-of-knowledge terminology, this is called witness extractability under discrete-logarithm hardness. The protocol becomes meaningful only in the font of DL hardness, which in turn guarantees that the protocol is secure.¹

2. THE PROBLEM OF DISTRIBUTED SCHNORR IDENTIFICATION

Suppose that some secret x with public counterpart y is distributed as x_1, \dots, x_n among n shareholders by means of Shamir’s (n, t) -sharing or any equivalent distributed key generation (DKG) protocol. That is, it shouldn’t matter if x is shared by a dealer or is implicitly defined in a fully decentralized fashion without being locally reconstructible; it only suffices that the output of the distribution scheme is indistinguishable from that of Shamir’s (n, t) -sharing. We are interested in a protocol that allows any $t^* \geq t$ (but no less) shareholders to collectively prove that they belong to the group represented by y in a zero-knowledge fashion. More accurately, involved shareholders should be able to prove that their individual shares combine to x by means of Shamir’s reconstruction formula (cf. Remark A.2) without revealing any information beyond that; evidently, this reduces to ordinary Schnorr identification for the edge case $t = 1$. Security against impersonation attacks should roughly mean that any efficient adversary controlling up to $t - 1$ shareholders remains unable to fabricate a proof that verifies with non-negligible chances. We obviously ask for the protocol to be as less interactive as possible; in the ideal case, involved shareholders should not need to engage in any communication rounds, performing instead single and independent requests to the verifier.

There is a seemingly trivial solution to this problem involving the public keys $y_i = g^{x_i}$, $1 \leq i \leq n$ of the shareholders. Let $Q \subset \{1, \dots, n\}$ be a coalition of shareholders that engage in an identification session with a verifier. The i -th party generates an ordinary Schnorr proof for its respective share, i.e.,

¹The most far-reaching result of this kind is that Schnorr identification is secure against active and concurrent impersonation attacks under “one-more” discrete logarithm (OMDL) hardness; cf. Theorem 2, [9].

$$(u_i, s_i), \quad s_i \leftarrow r_i + c_i x_i, \quad c_i \leftarrow H(u_i), \quad u_i \leftarrow g^{r_i}, \quad r_i \leftarrow_{\$} \mathbb{Z}_q$$

and sends it to the verifier. After aggregating the proof packets (u_i, s_i) , $i \in Q$, the verifier computes $c_i \leftarrow H(u_i)$, $i \in Q$ and accepts only if the relation

$$y = \prod_{i \in Q} y_i^{\lambda_i} \wedge g^{s_i} = y_i^{c_i} u_i, \quad \forall i \in Q \quad (2.1)$$

is satisfied, where $\{\lambda_i\}_{i \in Q}$ are the Lagrange interpolation coefficients for Q . We will refer to this as the *trivial approach*. That is, the verifier identifies the coalition only if the involved shareholders prove knowledge of their respective keys individually and their shares combine to the group public key in the sense of interpolation. Indeed, the latter is satisfied only if $|Q| \geq t$ (as desired), while the first is a common method for hardening distributed schemes against malleability, namely, requiring from the shareholders to present decoupled proofs for contextual statements that verify against their respective shares.² Note also that (2.1) is an overkill. If the public shares y_1, \dots, y_n are advertised as such anyway (i.e., the verifier registers them as the fixed sharing of y), the identifying condition reduces to

$$|Q| \geq t \wedge g^{s_i} = y_i^{c_i} u_i, \quad \forall i \in Q; \quad (2.2)$$

even if the threshold parameter does not become known along with the public shares, the verifier can trivially infer t from y_1, \dots, y_n and y and make use of it in subsequent protocol rounds.

Evidently, this protocol is secure if y_1, \dots, y_n is a (n, t) -sharing of y ; however, in the absence of side input on the verifier's side, the protocol is secure *only* under the latter assumption. Indeed, as illustrated by (2.2), the bulk of the work has been absorbed in the registration of the public shares during the setup phase; the verifier accepts only because the registered keys have already been trusted. Equivalently, the security analysis reduces trivially to the thread model for the initial setup. The problem is that, in real life, taking a trusted multi-key setup for granted is itself the opposite of trivial, since the possibilities regarding roles, trust assumptions and network operations compose a rather broad attack surface. In particular, since proof verification is entirely dependent on the registered public shares, the protocol remains susceptible to all sorts of malicious acting during the setup phase. For example, a proxy that is responsible for advertising the public shares may tamper with a certain subset of them, causing denial-of-service whenever the respective shareholders engage in a session. Or, worse, the proxy colludes with a malicious registry delivering the group public key and mounts the following rogue-key attack. The proxy retains y_1, \dots, y_{t-1} , replaces y_t with some y_t^* for which it knows the discrete logarithm and instructs the registry to deliver

$$y^* = (y_t^*)^{\lambda_t} \prod_{i=1}^{t-1} y_i^{\lambda_i}$$

in place of y , where $\{\lambda_i\}_i$ are the Lagrange interpolation coefficients for $\{1, \dots, t\}$. By reverse-engineering Shamir reconstruction, the proxy tunes y_i^* , $t < i \leq n$ so

²Compare the decoupled Chaum-Pedersen proofs in a threshold decryption scheme, used to ensure validity of the partial decryptors with respect to the ciphertext.

that $y_1, \dots, y_{t-1}, y_t^*, \dots, y_n^*$ becomes a (n, t) -sharing of y^* and sends it to the verifier; in the sequel, it can jump in and impersonate the t -th shareholder whenever the shareholders $1, \dots, t-1$ engage in a session. While ad hoc threat models can obviously be formulated for specific use cases (e.g., for a restricted pool of shareholders that undergo a well defined certification process), this discussion implies that an encompassing threat model is impossible to pin down. In the terminology of Bellare and Neven [12], the protocol has no chance to be provably secure in the *plain public-key* model.

One can alleviate the situation by working in the *knowledge of secret key* (KOSK) model and design a zero-knowledge based registration process, capable of detecting rogue-key attempts by means of purely cryptographic techniques. In this case, a security analysis becomes more tractable because a threat model for the initial setup can be formulated in terms of well defined malicious acts (even though dependence on the setup may introduce unexpected complexities to a formal security proof). However, an operation of this kind is not obvious to scale, since it should ensure that the public shares are compatible, and its repetition can become cumbersome if the shareholders are not static. More importantly, the KOSK assumption narrows the applicability of the protocol severely, since any verifiable registration process is an infrastructure plugin which real-world vendors may be unable or unwilling to interoperate with.

These are standard arguments speaking in favour of secure multi-party protocols that are agnostic with respect to their setup internals. For threshold identification, this requirement is further hardened as *non-reliance on the public shares*, meaning that the verification operation should only involve the group public key while the protocol remains secure in the plain model. In fact, this is a sine qua non if the desired protocol is to be applicable in a high-stressed and flexible decentralized setting (cf. Section 2.1).

2.1. Distributed provers in decentralized networks. We contended above that non-reliance on the public shares is a necessary condition for pragmatic security in the plain public-key model, where no concrete certification process for the shareholders' keys can be assumed to be available or known. Existence of such processes usually indicates a relatively limited number of shareholders that either belong to an organization or participate in a supposedly decentralized network, where group creation is subject to uniform rules and monitored. In such settings, the role of verifier is usually reserved for a bunch of authorized servers or dedicated nodes, responsible for managing some protected resource or performing specific acts in the context of a broader protocol. The trivial approach with an ad hoc initial setup, supported by the relevant infrastructure and accompanying attack model, can be a decent solution for these cases, at least insofar as the group of shareholders remains relatively static and general interoperability is not a hard requirement.

We are interested in use cases beyond that, where even maintaining the public shares, let alone trusting them, can cause severe headache. As a direct consequence, note that non-reliance on the public shares yields an almost plug-and-play solution for already advertised identities. Specifically, if the verification operation involves only the group public key, existing Schnorr verifiers need only extend their software in order to support the distributed case, without adding new infrastructure for the registration of newly advertised shares. Pre-existing keys can be silently shared at the verifiers' complete negligence.

We address distributed identification in a fairly unconditioned context, with as few assumptions as possible beyond a generic synchronous network with reliable packet delivery. The only assumption about Shamir-backed group creation should be that the public key of the combined prover be advertised in some “trusted” fashion; no group relevant infrastructure is assumed to be available after that phase, e.g., proxies, registries for individual member identities, storages for pre-processed quantities, aggregators interacting with involved provers, or even traceable communication channels among them. A group might be created on the fly and its public key be the only remnant of that ceremony (except for the secrets held by its members in private); after dissolution of the setup channels, its members may have no way to coordinate or identify with each other, be unwilling to do so or be compelled to act independently. This is even more important if the desired primitive is to be applicable as distributed proof-of-secret in a secure multiparty computation; in this case, combined keys do not usually outlive protocol rounds and any relevant infrastructure is a costly overkill (if not at all prohibited).

Absence of group infrastructure implies that the shareholders should in general be able to act without coordination during the proving phase. In particular, they should not even need to communicate, meaning that the protocol has to be non-interactive on the provers’ side. This introduces a high degree of asynchronicity which in turn imposes constraints on the interaction of the combined prover with the verifier. In particular, a verifier’s response to the j -th shareholder should not be dependent on any request from the i -th shareholder for $i \neq j$. Equivalently, the proving session should decouple into independent interactions between the verifier and the individual shareholders; after aggregating the decoupled results of these interactions, the verifier combines them and proceeds to the final verification step. Ideally, the decoupled interactions should themselves be non-interactive, i.e., consist of a single communication step in which individual shareholders send their respective proof packets to the verifier.

Note that the trivial approach satisfies these requirements automatically. Indeed, it presupposes no coordination or infrastructure to be available during the proving phase; infrastructure is only needed during the setup in order to accommodate the verifiers’ trust on the public shares. To see why this is still inadequate, consider a *truly* decentralized network where qualified verifiers can be user devices or application servers alike. That is, verification may occur massively on a multitude of devices that need not be aware of the individual shareholders’ identities for every group they come across in the network. It simply cannot be expected from a mobile or other IoT device to maintain the public shares of every group it might want to verify, let alone engage in a certification process in order to register them. Note also that fetching the public shares from remote registries on demand is not generally an option; even if maintaining these registries were possible, trusting them is a highly non-trivial assumption (comparable to KOSK) while depending on them would significantly affect network latency and robustness.

Further efficiency considerations are here in order, since decentralized verifiers do not necessarily possess the capabilities of a fine-tuned application server. It is well known that the wireless transmission of a single bit consumes more power than a 64-bit (or 32-bit) instruction by orders of magnitude. Consequently, reduced throughput is usually the dominating factor with respect to overall energy

consumption; from the individual node's perspective, this amounts to longer battery life if it runs on a wireless device. Reduced throughput is even more important in order to avoid network congestion; indeed, if the network is truly decentralized, proof packets have to be massively transmitted through a bunch of dedicated relay nodes in order to be verified at arbitrary locations. The problem remains to design a non-interactive threshold protocol with reduced throughput that is non-reliant on public shares, infrastructure-free on both sides, secure in the plain public-key model, and as such applicable in truly decentralized settings.

2.2. Non-working solutions. We informally discuss two plausible approaches to the problem of distributed Schnorr identification and show how they fail to meet the requirements of Section 2.1. Both are within the plain public-key model.

2.2.1. Uniform commitment approach. Recall that Schnorr signature is Schnorr identification with a message baked into the hash function. This suggests an approach similar to the threshold signature scheme introduced by Stinson and Stroh [6]. This requires no infrastructure extensions on the verifier's side, since verification involves only the group public key. Let x_1, \dots, x_n be the fixed (n, t) -sharing of some secret x with public counterpart y . A coalition of shareholders $Q \subset \{1, \dots, n\}$, $|Q| \geq t$, execute an equivalent DKG protocol to generate

$$r = \sum_{i \in Q} \lambda_i r_i,$$

where $\{\lambda_i\}_{i \in Q}$ are the interpolation coefficients for Q , holding the respective r_i 's in private. r need not be reconstructible at any single location but involved parties should be able to verifiably infer the distributed commitment $u = g^r$ at the end of the process. Next, the involved parties generate Schnorr proofs for their keys using the respective commitment share and uniform challenge $H(u)$, i.e.,

$$(u_i, s_i), \quad s_i \leftarrow r_i + cx_i, \quad u_i \leftarrow g^{r_i}, \quad c \leftarrow H(u),$$

and send them to the verifier. After aggregating (u_i, s_i) , $i \in Q$, the verifier computes

$$s \leftarrow \sum_{i \in Q} \lambda_i s_i, \quad c \leftarrow H(u), \quad u \leftarrow \prod_{i \in Q} u_i^{\lambda_i}$$

and accepts only if the relation $g^s = y^c u$ is satisfied. Indeed, since $|Q| \geq t$, twofold application of Shamir's reconstruction formula yields

$$g^s = \prod_{i \in Q} (g^{s_i})^{\lambda_i} = \left(\prod_{i \in Q} (g^{x_i})^c \right)^{\lambda_i} \prod_{i \in Q} (g^{r_i})^{\lambda_i} = \left(\prod_{i \in Q} y_i^{\lambda_i} \right)^c \prod_{i \in Q} u_i^{\lambda_i} = y^c u.$$

Verification requires $|Q| + 2$ exponentiations and the inbound throughput grows linearly with respect to $|Q|$; for example, taking \mathbb{G} to be the elliptic curve P256 and assuming that 128-bit challenges is sufficient to work with, the total throughput is $512|Q|$ bits per protocol round. If accommodated, we can offload the computation of s and u to a semi-trusted combiner, who acts as group proxy by forwarding (u, s) to the verifier; in this case, the verifier coincides with that of the non-distributed

case and its performance is not dependent on the number of involved shareholders. This is usually the case for Schnorr threshold signatures.

We will refer to the above protocol as the *uniform commitment* approach. Its security reduces most probably to that of the non-distributed case by carefully adapting the arguments of [6] to the identification setting. It relies on the fact that this “is” ordinary Schnorr identification if we treat the involved shareholders as a single entity; indeed, the combined proof (u, s) lies in the domain of proofs that could have been generated for x directly. This commutation is crucial for a threshold signature scheme, where the final object may need to remain available in the long run and consequently use minimum space, but not for an identification protocol, where no persistent data need be generated. In fact, it countervenes the requirement of reduced interactivens, as it is achieved through a highly interactive computation for generating g^r . This evidently introduces significant network overhead upon every protocol round and presupposes that some infrastructure is constantly available to the group of shareholders. Consequently, despite the efficiency of its verifier, the uniform commitment approach does not meet the requirements of Section 2.1 on the provers’ side. The state-of-the-art for managing the network overhead of Schnorr threshold signatures is the FROST protocol introduced by Komlo and Goldberg [3] and its variants. This reduces the number of communication rounds to one or two by means of a semi-trusted aggregator, depending on whether a preprocessing stage takes place during the initial setup; however, it comes at the cost of more infrastructure (and consequently more trust assumptions) which makes it even more prohibiting for a flexible identification setting.

2.2.2. Completely decoupled approach. This is a variance of the trivial approach, where the public shares need not be advertized and trusted during the setup phase. Instead, they are attached to the Schnorr proofs that are generated by the respective shareholders. Given a (n, t) -sharing x_1, \dots, x_n of some secret x with public counterpart $y = g^x$, a coalition of shareholders $Q \subset \{1, \dots, n\}$ generate

$$(y_i, u_i, s_i), \quad y_i \leftarrow g^{x_i}, \quad s_i \leftarrow r_i + c_i x_i, \quad c_i \leftarrow H(u_i), \quad u_i \leftarrow g^{r_i}, \quad r_i \leftarrow_{\$} \mathbb{Z}_q$$

respectively and send them to the verifier. After aggregating (y_i, u_i, s_i) , $i \in Q$, the verifier computes $c_i \leftarrow H(u_i)$ and accepts only if condition (2.1) is satisfied. We will refer to this protocol as the *completely decoupled* approach. Intuitively, its security reduces to the proof-of-knowledge property of the Schnorr protocol. Every shareholder “should know” the discrete logarithm of the attached share because the respective proof verifies; since the attached shares combine to the group public key, their respective logarithms comprise a sharing of x and the shareholders belong to the group represented by y .

The protocol is non-reliant on public shares and infrastructure-free on both sides. However, this is achieved at the cost of $3|Q|$ exponentiations per verification and at least 33% increased throughout as compared to the trivial approach. While this regression might be insignificant for a central identifying server with even several thousands of users, it can badly affect efficiency under the energy intensive and highly congested network conditions of true decentralization (cf. Section 2.1).

Moreover, the security argument is not as strong as it might seem. The Schnorr proof-of-knowledge property is not concretely defined in terms of an attack game

and it basically means extractability under discrete logarithm hardness. Consequently, the security of the threshold protocol cannot be formally reduced to it in the strict sense; by saying that it “reduces to” it, what we actually mean is that some relevant extraction procedure should itself be reproducible in the context of a black-box threshold attack. This does not seem to be possible. Consider an adversary \mathcal{A} who knows $\{x_i\}_{i \in J}$ for some $J \subset \{1, \dots, n\}$ with $|J| = t - 1$ and wants to impersonate $Q = J \cup \{j\}$ with $j \notin J$ against an interactive verifier. One possible attack pattern could be using the shares that are known to the adversary directly. That is, for every $i \in J$, \mathcal{A} sends

$$(y_i, u_i), \quad y_i \leftarrow g^{x_i}, \quad u_i \leftarrow g^{r_i}, \quad r_i \leftarrow_{\$} \mathbb{Z}_q$$

to the verifier, who samples $c_i \leftarrow_{\$} \mathbb{Z}_q$ for u_i and sends it to \mathcal{A} ; in turn, \mathcal{A} responds with $s_i \leftarrow r_i + c_i x_i$. Next, \mathcal{A} fabricates u_j according to some black-box strategy and sends (y_j, u_j) , $y_j \leftarrow g^{x_j}$ to the verifier, who samples $c_j \leftarrow_{\$} \mathbb{Z}_q$ for u_j and sends it to the adversary. Finally, \mathcal{A} fabricates s_j according to some black-box strategy and sends it to the verifier. Evidently, extraction of x_j is here possible by rewinding \mathcal{A} exactly before sending (y_j, u_j) , provided that it succeeds in generating a valid transcript. This is ordinary extraction embedded into a decoupled threshold attack pattern. However, a black-box threshold attack is far from confined to this pattern. Since y admits many possible sharings, \mathcal{A} could possibly fabricate attachments that do not coincide with the fixed public shares and still combine to the group public key; in this case, rewinding \mathcal{A} at any location would fail to yield x_j . We would need the decoupled proofs to be further junct into a single operation in order for the extraction of x_j to be possible by rewinding.

3. ONE-ROUND SCHNORR THRESHOLD IDENTIFICATION

Throughout this section we fix a cyclic group $\mathbb{G} = \langle g \rangle$ of order q where the discrete-logarithm problem is hard and a secure hash function $H : \mathbb{G} \rightarrow \mathbb{Z}_q$. The Shamir’s (n, t) -sharing is denoted by D (cf. Section A), with the particular values of its parameters inferred always from context. Given integers $1 \leq t \leq n < q$ TODO

3.1. The ORST protocol.

Protocol 3.1. (ORST $_{\mathbb{G}, n, t}$ – *One-Round Schnorr (n, t) -Threshold Identification*)

- *Setup.* Some uniformly random secret is distributed among a group of n shareholders by means of Shamir’s (n, t) -sharing, i.e.,

$$(x_1, \dots, x_n; y) \leftarrow D(x), \quad x \leftarrow_{\$} \mathbb{Z}_q$$

The threshold parameter and the public shares may remain private to the group, who needs only advertise the combined public key y .

- *Proving phase.* A coalition of shareholders $Q \subset \{1, \dots, n\}$ respectively send

$$(u_i, s_i), \quad s_i \leftarrow r_i + c_i x_i, \quad c_i \leftarrow H(u_i), \quad u_i \leftarrow g^{r_i}, \quad r_i \leftarrow_{\$} \mathbb{Z}_q$$

to the verifier and hold r_i in private, $i \in Q$.

◦ *Verification.* After aggregating (u_i, s_i) , $i \in Q$ the verifier accepts only if

$$\exp\left(g, \sum_{i \in Q} \mu_i s_i\right) = y^{\bar{c}} \prod_{i \in Q} u_i^{\mu_i}, \quad (3.2)$$

where

$$\bar{c} \leftarrow \prod_{i \in Q} c_i, \quad \mu_i \leftarrow \lambda_i \prod_{j \in Q \setminus \{i\}} c_j, \quad c_i \leftarrow H(u_i) \quad (3.3)$$

and $\{\lambda_i\}_{i \in Q}$ are the interpolation coefficients for Q .

Remark 3.4. (Validity condition) By equating exponents, (3.2) translates to

$$\sum_{i \in Q} \mu_i (s_i - r_i) = \bar{c} \cdot x, \quad (3.5)$$

where $r_i \equiv \log u_i$. Observe that this contains no reference to shares.

Remark 3.6. (Correctness) Let $y_i = g^{x_i}$, $1 \leq i \leq n$ be the public shares. The left-hand side of (3.2) transforms as

$$\prod_{i \in Q} (g^{s_i})^{\mu_i} = \prod_{i \in Q} (g^{x_i})^{c_i \mu_i} (g^{r_i})^{\mu_i} = \prod_{i \in Q} (g^{x_i})^{\lambda_i \bar{c}} \prod_{i \in Q} (g^{r_i})^{\mu_i} = \left(\prod_{i \in Q} y_i^{\lambda_i} \right)^{\bar{c}} \prod_{i \in Q} u_i^{\mu_i},$$

which yields the right-hand side if (and only if) $|Q| \geq t$. In other words, the proper execution of the protocol verifies only if at least t shareholders are involved.

Remark 3.7. (Interactive aspect) We may regard H as approximating a random oracle insofar as predicting any of its output bits remains infeasible; in the limit, involved shareholders submit their commitments to an external oracle that is accessible to the verifier. Or, equivalently, they send their commitments directly to the verifier, who lazily samples challenges by maintaining a lookup table **Map**, so that the protocol becomes interactive. Specifically, the i -th shareholder samples $r_i \leftarrow_{\$} \mathbb{Z}_q$ and sends the commitment $u_i \leftarrow g^{r_i}$ to the verifier, who samples $c_i \leftarrow_{\$} \mathbb{Z}_q$ if $u \notin \text{Dom}(\text{Map})$, caches $\text{Map}[u_i] \leftarrow c_i$ and responds with c_i ; otherwise, it responds with $c_i \leftarrow \text{Map}[u_i]$. In turn, the i -th shareholder responds with $s_i \leftarrow r_i + c_i x_i$. After aggregating s_i , $i \in Q$, the verifier checks that condition (3.2) is satisfied using the previously cached challenges $c_i = \text{Map}[u_i]$, $i \in Q$.

The sharing scheme of the setup may be replaced with any DKG whose output distribution is indistinguishable from that of Shamir's secret sharing. That is, it shouldn't matter if x has been shared by a dealer or it cannot be reconstructed at a single location; it only suffices that the correct public key $y = g^x$ be advertised at the end of the process. Note that DKG's usually assume a bound $0 \leq l < t$ on the number of corrupt parties in order for the distribution to be secure. Since a threshold threat model allows the corruption of up to $t - 1$ parties, plugging a DKG is indifferent only if the corruption takes place after the completion of key distribution. This is in accordance with the requirement that ORST should be agnostic with respect to its setup internals, meaning that a separate security analysis should be dedicated per use case to it. TODO

During the proving phase, the involved shareholders generate ordinary Schnorr proofs for their respective shares and send them to the verifier asynchronously.

No agreement or coordination is assumed between them. Order coordination or some session timeout may be enforced by the verifier for the needs of a specific use case, but this should not affect the threat model essentially because the local computations remain decoupled. The protocol is non-interactive in a strict sense (“one-round”), since it requires no communication rounds among the shareholders. Specifically, ORST meets the requirements of Section 2.1 on the provers’ side.

Verification efficiency is comparable to that of the uniform challenge approach (without proxy), even if at the cost of more numerical operations. It requires $|Q|+2$ exponentiations while its throughput grows linearly with respect to the number of involved shareholders; for example, if \mathbb{G} is the elliptic curve P256 and assuming that 128-bit challenges is sufficient to work with, the inbound throughput is $512|Q|$ bits per protocol round. Since the operation per se involves the group public key alone, no maintenance of the public shares is needed and the protocol meets the requirements of Section 2.1 on the verifier’s side. Existing Schnorr verifiers need only extend their software in order to support the distributed case, without adding new infrastructure for the certification and registration of public shares.

Definition 3.8. The *weights* of a collection $\{u_i\}_{i \in Q} \subset \mathbb{G}$, $Q \subset \{1, \dots, n\}$, are the scalars $\{\mu_i\}_{i \in Q}$ defined as

$$\mu_i = \lambda_i \prod_{j \in Q \setminus \{i\}} c_j \quad (3.9)$$

where $c_i \equiv H(u_i)$, $i \in Q$ and $\{\lambda_i\}_{i \in Q}$ are the interpolation coefficients for Q . Under the same notation, the *normalizer* of $\{u_i\}_{i \in Q}$ is the scalar

$$\bar{c} = \prod_{i \in Q} c_i \quad (3.10)$$

That is, ORST-verification juncts the decoupled Schnorr proofs through the weights and normalizer of the points $\{u_i\}_{i \in Q}$, which the shareholders respectively commit to. We have already used (cf. Remark 3.6) the following fundamental remark.

Remark 3.11. (Relation between weights and the normalizer) Let $\{\mu_i\}_{i \in Q}$, \bar{c} be the weights and the normalizer of a collection $\{u_i\}_{i \in Q}$ respectively. Then

$$c_i \mu_i = \lambda_i \bar{c}, \quad \forall i \in Q \quad (3.12)$$

Definition 3.13. By *proof* is meant a collection $\sigma = \{(u_i, s_i)\}_{i \in Q}$ of the form

$$(u_i, s_i) \in \mathbb{G} \times \mathbb{Z}_q, \quad \forall i \in Q \quad \wedge \quad Q \subset \{1, \dots, n\}.$$

Given σ as above, we refer to $\{u_i\}_{i \in Q}$ as its *commitments* and say that σ is *based on* them. We also refer to $\{s_i\}_{i \in Q}$ as the *responses*. The *weights* resp. *normalizer of* σ are the weights resp. normalizer of its commitments. The *proof space* is the set of proofs and will be denoted by Σ . We also denote by

$$\Sigma[\{u_i\}_{i \in Q}] = \{\sigma \in \Sigma : \exists \{s_i\}_{i \in Q} \text{ such that } \sigma = \{(u_i, s_i)\}_{i \in Q}\} \quad (3.14)$$

the subspace of proofs that are based on $\{u_i\}_{i \in Q}$. Evidently, proofs in this subspace share the same weights and normalizer.

Note that the above proof definition does not impose constraints on the relation between the commitments and the responses. Specifically, it does not confine to properly generated transcripts; it models the verifier's view in the wild and includes anything that malicious provers can possibly fabricate. The proper execution of the protocol is modelled as follows.

Definition 3.15. $\{(u_i, s_i)\}_{i \in Q}$ is called *canonical with respect to* $(x_1, \dots, x_n; y) \in D_x$ if its components are ordinary Schnorr proofs for the respective shares, i.e.,

$$s_i = r_i + c_i x_i, \quad c_i \equiv H(u_i), \quad r_i \equiv \log u_i, \quad \forall i \in Q \quad (3.16)$$

By the “unique responses” property of the ordinary Schnorr protocol, $\Sigma[\{u_i\}_{i \in Q}]$ contains a unique canonical representative with respect to every fixed sharing.

Definition 3.17. The *one-round Schnorr* (n, t) -threshold verifier is the deterministic algorithm $\mathcal{V} : \mathbb{G} \times \Sigma \rightarrow \{\text{true}, \text{false}\}$ invoked as

$$(y, \{(u_i, s_i)\}_{i \in Q}) \mapsto \exp\left(g, \sum_{i \in Q} \mu_i s_i\right) = y^{\bar{c}} \prod_{i \in Q} u_i^{\mu_i},$$

where $\{\mu_i\}_{i \in Q}$, \bar{c} are the weights and the normalizer of $\{(u_i, s_i)\}_{i \in Q}$ respectively.

Definition 3.18. A proof $\sigma \in \Sigma$ is called *valid against* $y \in \mathbb{G}$ if $\mathcal{V}(y, \sigma) = \text{true}$.

Remark 3.19. ORST correctness (cf. Remark 3.6) means that a proof canonical with respect to $(x_1, \dots, x_n; y)$ is valid against y if and only if $|Q| \geq t$.

The subspace of proofs based on $\{u_i\}_{i \in Q}$ and valid against y is denoted by

$$\Sigma_y[\{u_i\}_{i \in Q}] = \{\sigma \in \Sigma[\{u_i\}_{i \in Q}] : \mathcal{V}(y, \sigma) = \text{true}\}. \quad (3.20)$$

If $|Q| \geq t$, then this space contains a unique canonical representative with respect to every fixed sharing of $x = \log y$.

Remark 3.21. In the non-distributed case $t = 1$, the space $\Sigma_y[\{u_i\}_{i \in Q}]$ is a singleton due to the “unique responses” property of the Schnorr protocol. Specifically, for every commitment u there exists a unique response s such that the proof (u, s) is valid against $y = g^x$, namely $s = r + H(u)x$ with $r \equiv \log u$. In general, however, this space inflates dramatically at a rate dependent on $|Q| \geq t$ (cf. Remark 3.36). This implies an abundance of valid proofs that are non-canonical, a fact that every meaningful security notion should be able to trace.

Proposition 3.22. (*Validity conditions*) Let $x \in \mathbb{Z}_q$, $y = g^x$ and $\sigma = \{(u_i, s_i)\}_{i \in Q}$. Then, with overwhelming probability, σ is valid against y if and only if

$$\frac{1}{\bar{c}} \sum_{i \in Q} \mu_i (s_i - r_i) = x, \quad r_i \equiv \log u_i, \quad (3.23)$$

where $\{\mu_i\}_{i \in Q}$, \bar{c} are its weights and normalizer respectively or, equivalently,

$$x = \sum_{i \in Q} \lambda_i \frac{s_i - r_i}{c_i}, \quad (3.24)$$

where $c_i \equiv H(u_i)$ and $\{\lambda_i\}_{i \in Q}$ are the interpolation coefficients for Q .

Proof. Since H is a secure hash function, $\bar{c} \neq 0$ with overwhelming probability in the bitlength of the group order. In this case, the verification condition (3.5) becomes (3.23). It further reformulates to (3.24) by application of (3.12). \square

This statement holds true irrespective of size (including the case $|Q| < t$), establishing a relation between x and any proof that happens to be valid against y without reference to sharings. Both (3.23) and (3.24) generalize the relation of a secret to a valid non-distributed proof; it should be stressed however that (3.24) does not imply $x_i = (r_i - s_i)/c_i$ and can hold true without σ being canonical. Indeed (cf. Remark 3.29), valid proofs do not confine to those properly generated in ORST and a relevant security notion should ensure that such proofs are infeasible to compute without knowledge of at least t shares (cf. Section 4). The exact meaning of (3.24) is clarified in Proposition 3.28. A proof should evidently not verify against different public keys. In fact, this key almost over exists and it is unique.

Proposition 3.25. *(Almost every proof is valid against some unique public key) Given commitments $\{u_i\}_{i \in Q} \subset \mathbb{G}$, then, with overwhelming probability, every $\sigma \in \Sigma$ based on them is valid against some unique $y \in \mathbb{G}$.*

Proof. Let $\{\mu_i\}_{i \in Q}$, \bar{c} be the weights and the normalizer of $\{u_i\}_{i \in Q}$, which are common to all proofs based on them. Since H is a secure hash function, $\bar{c} \neq 0$ with overwhelming probability in the bitlength of the group order. Consequently, given any proof of the form $\sigma = \{(u_i, s_i)\}_{i \in Q}$, we can define

$$x = \frac{1}{\bar{c}} \sum_{i \in Q} \mu_i (s_i - r_i),$$

where $r_i \equiv \log u_i$. By Proposition 3.22, σ is valid against $y = g^x$ solely. \square

Corollary 3.26. *Given $\{u_i\}_{i \in Q} \subset \mathbb{G}$, then, with overwhelming probability,*

$$\Sigma[\{u_i\}_{i \in Q}] = \bigcup_{y \in \mathbb{G}} \Sigma_y[\{u_i\}_{i \in Q}] \quad (3.27)$$

where the right-hand side is a partition, i.e.,

$$\Sigma_y[\{u_i\}_{i \in Q}] \cap \Sigma_{y^*}[\{u_i\}_{i \in Q}] = \emptyset \quad \text{for } y \neq y^*.$$

Proof. Direct consequence of Proposition 3.25 and the definitions. \square

Proposition 3.28. *(Almost every proof of size $|Q| = t$ is canonical with respect to some unique sharing) Given commitments $\{u_i\}_{i \in Q} \subset \mathbb{G}$ with $|Q| = t$, then, with overwhelming probability, every $\sigma \in \Sigma$ based on them is canonical with respect to some unique sharing $(x_1, \dots, x_n; y) \in D_x$ of some unique $x \in \mathbb{Z}_q$.*

Proof. Let $\{\mu_i\}_{i \in Q}$ be the weights of $\{u_i\}_{i \in Q}$ and denote $c_i \equiv H(u_i)$. Note that these parameters are common to all proofs based on the given commitments. Since H is a secure hash function, $c_i \neq 0$, $\forall i \in Q$ with overwhelming probability in the bitlength of the group order. Consequently, given $\sigma = \{(u_i, s_i)\}_{i \in Q}$, we can define

$$x_i = \frac{s_i - r_i}{c_i}, \quad i \in Q$$

where $r_i \equiv \log u_i$. By virtue of interpolation, since $|Q| = t$, the scalars $\{x_i\}_{i \in Q}$ uniquely determine a (n, t) -sharing $(x_1, \dots, x_n; y)$ of some $x \in \mathbb{Z}_q$, namely

$$x = \sum_{i \in Q} \lambda_i x_i.$$

Since $s_i = r_i + c_i x_i, \forall i \in Q$, the proof is canonical with respect to this sharing and thus valid against $y = g^x$. The claim follows because a proof can be valid with respect to at most one public key (cf. Prop. 3.25). \square

Remark 3.29. (Validity and degrees of freedom) Fix commitments $\{u_i\}_{i \in Q} \subset \mathbb{G}$. Choose $J \subset Q$ with $|J| = |Q| - 1$, pick arbitrary $\{s_i\}_{i \in J} \subset \mathbb{Z}_q$ and set

$$s_j = r_j + \frac{1}{\mu_j} \left(\bar{c} \cdot x - \sum_{i \in J} \mu_i (s_i - r_i) \right)$$

where $\{j\} = Q \setminus J$ and $r_i \equiv \log u_i$. By Proposition 3.22 (cf. (3.23)), this process exhausts the proofs that are based on $\{u_i\}_{i \in Q}$ and valid against $y \equiv g^x$ bijectively. That is, the subspace $\Sigma_y[\{u_i\}_{i \in Q}]$ has $|Q| - 1$ degrees of freedom in the sense that every point of it is uniquely determined by the choice of $\{s_i\}_{i \in J}$ and $J \subset Q$ (for $|Q| = 1$ this is the “unique responses” property of ordinary Schnorr identification). Note that this process has no obvious utility for an impersonating adversary, since the adversary would still need to eliminate x in order to compute s_j .

Proposition 3.30. (*Validity against sharing*) Let $x \in \mathbb{Z}_q$ and $\sigma = \{(u_i, s_i)\}_{i \in Q}$. Given $(x_1, \dots, x_n; y) \in D_x$, then σ is valid against y if and only if

$$\sum_{i \in Q} \mu_i (s_i - r_i - c_i x_i) = \bar{c} \cdot \left(x - \sum_{i \in Q} \lambda_i x_i \right), \quad (3.31)$$

Proof. Follows from the verification condition (3.5) by subtracting

$$\sum_{i \in Q} \mu_i c_i x_i$$

on both sides and applying (3.12) on the right. \square

Remark 3.32. (Valid proofs with $|Q| < t$) We know that if a proof of size $< t$ is valid, then it cannot be canonical (cf. Remark 3.6). Such proofs exist in abundance (cf. Remark 3.29) and Prop. 3.30 indicates a way to generate them programmatically. Fix $\{u_i\}_{i \in Q} \subset \mathbb{G}$ and set $s_i = r_i + c_i x_i$ where $r_i \equiv \log u_i$, $c_i \equiv H(u_i)$. Pick $S \subset Q$ with $S \neq \emptyset$ and define $s_i^* = s_i$ if $i \notin S$, otherwise

$$s_i^* = s_i + \delta_i, \quad \delta_i \equiv \frac{1}{\mu_i} \cdot \frac{\bar{c}}{|S|} \cdot \left(x - \sum_{j \in Q} \lambda_j x_j \right),$$

where \bar{c} is the normalizer of $\{u_i\}_{i \in Q}$. Consider the proof $\sigma = \{(u_i, s_i)\}_Q$. If $|Q| \geq t$, then $\delta_i = 0, \forall i \in S$ so that σ is canonical and thus automatically valid. If $|Q| < t$, then σ still remains valid because it satisfies (3.31) by design. Note however that this construction has no obvious utility for a malicious coalition $Q \subset \{1, \dots, n\}$ with $|Q| < t$ since the colluders (or, equivalently, an impersonating adversary who knows $\{x_i\}_{i \in Q}$) would still need to know x in order to adjust the δ_i ’s appropriately.

Corollary 3.33. (*Validity against sharing, $|Q| \geq t$*) Let $x \in \mathbb{Z}_q$, $\sigma = \{(u_i, s_i)\}_{i \in Q}$ with $|Q| = t$. Given $(x_1, \dots, x_n; y) \in D_x$, then σ is valid against y if and only if

$$\sum_{i \in Q} \mu_i s_i = \sum_{i \in Q} \mu_i (r_i + c_i x_i), \quad (3.34)$$

where $r_i \equiv \log u_i$, $c_i \equiv H(u_i)$ and $\{\mu_i\}_{i \in Q}$ are the weights of σ .

Proof. The right-hand side of (3.31) vanishes exactly if $|Q| \geq t$. \square

Remark 3.35. (*$t - 1$ degrees of freedom*) Fix $\{u_i\}_{i \in Q}$ with $|Q| = t$. Choose $J \subset Q$ with $|J| = t - 1$, pick arbitrary $\{\delta_i\}_{i \in J} \subset \mathbb{Z}_q$ and set

$$s_j = r_j + c_j x_j - \frac{1}{\mu_j} \sum_{i \in J} \mu_i (s_i - r_i - c_i x_i),$$

where $\{j\} = Q \setminus J$ and $r_i \equiv \log u_i$, $c_i \equiv H(u_i)$. By Corollary (3.33), this process exhausts the proofs that are based on $\{u_i\}_{i \in Q}$ and valid against y bijectively. This is a special case of Remark 3.29 with reference to a fixed sharing. Note however that it has no obvious utility for an adversary who knows $\{x_i\}_{i \in J}$ but not x_j .

Remark 3.36. (*Valid proofs with $|Q| \geq t$*) This can be seen from a more general perspective. Consider a canonical proof $\{(u_i, s_i)\}_{i \in Q}$ with $|Q| \geq t$ and let $\{\mu_i\}_{i \in Q}$ be its weights. Pick $j \in Q$ and arbitrary scalars δ_i , $i \in Q \setminus \{j\}$ such that at least one of them is non-zero. If we define

$$\delta_j = -\frac{1}{\mu_j} \sum_{i \neq j} \mu_i \delta_i,$$

then the proof $\{(u_i, s_i + \delta_i)\}_{i \in Q}$ satisfies condition (3.34) by design. Evidently, this process exhausts the valid proofs that are based on $\{u_i\}_{i \in Q}$.

3.2. Security considerations. Since ORST decouples into concurrent runs of ordinary Schnorr identification, any security precaution for the latter applies also in the threshold case. Most notably, security against side-channel attacks cannot be attained if the random number generator leaks linear relations between the discrete logarithms of distinct commitments. We here focus on high-level attacks that arise particularly in the distributed setting.

3.2.1. Generalized replay attack. ORST is as much susceptible to replay attacks as the non-distributed Schnorr protocol. That is, a man-in-the-middle can passively capture and anytime replay the packets sent by a coalition of shareholders in order to impersonate them. Things seem worse because the attacker can transform the intercepted proof almost arbitrarily, provided that the commitments are kept fixed. Let $\{(u_i, s_i)\}_{i \in Q}$, $|Q| \geq t$ be a valid proof with weights $\{\mu_i\}_{i \in Q}$. For fixed $j \in Q$, denote $J = Q \setminus \{j\}$, pick $\{\delta_i\}_{i \in J} \subset \mathbb{Z}_q$ and define

$$\delta_j = -\frac{1}{\mu_j} \sum_{i \in J} \mu_i \delta_i.$$

Then $\{(u_i, s_i^*)\}_{i \in Q}$ with $s_i^* = s_i + \delta_i$ remains valid because

$$\sum_{i \in Q} \mu_i s_i^* = \sum_{i \in Q} \mu_i s_i + \left(\mu_j \delta_j + \sum_{i \in J} \mu_i \delta_i \right) = \sum_{i \in Q} \mu_i s_i$$

(cf. Cor. 3.33). This essentially exploits the structure described in Remark 3.36. The security of the protocol should ensure that an attacker controlling up to $t - 1$ shareholders remains unable to fabricate $\{(u_i, s_i^*)\}_{i \in Q}$ without having first intercepted $\{(u_i, s_i)\}_{i \in Q}$; still, in case of interception, the attacker has a least $t - 1$ degrees of freedom in fabricating a valid proof based on the same commitments. This generalizes the replay attack of the non-distributed case $t = 1$, where the intercepted proof can only be replayed as is. It should be nonetheless clear that generalized replay attacks are mitigated by the same method as in the non-distributed case, i.e., by baking a nonce into the hash function, capable of maintaining state between the verifier and the involved provers.

3.2.2. Non-accountability. TODO

3.2.3. *Key-recovery attack – interactive aspect.* This practically affects only the interactive version of the protocol (cf. Remark 3.7) under rather special circumstances, but sheds light on its inherent security structure. Similar to the non-distributed case, we employ this attack in order to derive a formal security proof by means of fork and extraction (cf. Section 5.5). Evidently, if a shareholder uses the same commitment in two canonical rounds of the interactive ORST (due to, say, broken random generator) then its secret share leaks trivially. This is the key-recovery attack of the non-distributed case embedded into the threshold setting. We will show that this holds true even for non-canonical executions given a malicious verifier who controls $t - 1$ shareholders.

Let $Q = J \cup \{j\}$ be a coalition of shareholders such that $|J| = t - 1$, $j \notin J$ and the cluster J is controlled by an adversary \mathcal{A}^* . We further assume that \mathcal{A}^* has compromised the identifying server. The shareholders engage in an identification session and generate a transcript $\{(u_i, c_i, s_i)\}_{i \in Q}$ that verifies, so that

$$\sum_{i \in Q} \mu_i (s_i - r_i - c_i x_i) = 0$$

holds true, where $r_i \equiv \log u_i$ and $\{\mu_i\}_{i \in Q}$ are the weights induced by the challenges $\{c_i\}_{i \in Q}$. (cf. Corollary 3.34). They next engage in a second session, where the j -th shareholder reuses the commitment u_j . After receiving u_j on the server's side, \mathcal{A}^* responds with a challenge $c_j^* \neq c_j$ and tunes the rest shareholders to reuse their commitments u_i , $i \in J$ from the previous session. To each of them, the server responds with the same challenge $c_i^* = c_i$ from the previous session. Subsequently, all shareholders compute honestly their respective responses s_i , $i \in Q$ and send them to the server; since the transcript $\{(u_i, c_i^*, s_i^*)\}_{i \in Q}$ verifies, we again have

$$\sum_{i \in Q} \mu_i^* (s_i - r_i - c_i^* x_i) = 0,$$

where $\{\mu_i^*\}_{i \in Q}$ are the weights induced by the challenges $\{c_i^*\}_{i \in Q}$. Note that, since $c_i = c_i^*$ for all $i \neq j$, by definition of weights (cf. Definition 3.8) we get $\mu_j = \mu_j^*$. Subtracting terms in the above relations and applying this remark, we get

$$\mu_j(s_j - s_j^* - (c_j - c_j^*)x_j) + \sum_{i \in J} (\mu_i(s_i - r_i - c_i x_i) - \mu_i^*(s_i^* - r_i - c_i x_i)) = 0.$$

This eliminates r_j , which is the only unknown to the adversary. Since $c_j - c_j^* \neq 0$, \mathcal{A}^* can use this relation to retrieve x_j in the form

$$x_j = \frac{s_j - s_j^*}{c_j - c_j^*} + \frac{1}{\mu_j} \sum_{i \in J} \left(\mu_i(s_i - r_i - c_i x_i) - \mu_i^*(s_i^* - r_i - c_i x_i) \right). \quad (3.37)$$

Note finally that since $|Q| = t$, the adversary can use the shares $\{x_i\}_{i \in Q}$ to compute the combined secret key by means of Shamir's reconstruction formula.

4. SECURITY NOTIONS

In this section we specify the threat model for the protocols of Section ?? . Formulations are given in terms of attack games within the asymptotic paradigm and negligibility is meant with respect to the bitlength of the group order. TODO The relevant formalism for discrete-log based settings is developed in Section ??.

4.1. Threat model. The current standard for identification protocols is security against active impersonation attacks. This refers to adversaries who not only passively intercept identification sessions before mounting their actual attack, but also engage in them actively by impersonating a legitimate verifier (e.g., by cloning a website in order to collect information about the user's credentials); after several such interactions, the adversary switches roles and attempts to identify itself against some verifier as a legitimate user. We will not pursue this path, focusing on adversaries who only intercept identification sessions and potentially exert some control on the network. Once the relation between threshold and ordinary Schnorr identification has been clarified from the provability viewpoint, security against active adversaries should be straightforward to establish using techniques similar to those of [9] due to the decoupled structure of the distributed protocol.

4.1.1. General threat model. In most general terms, the ORST threat model consists of an adversary \mathcal{A} who compromises the keys of up to $t - 1$ shareholders and attempts to impersonate some coalition $Q \subset \{1, \dots, n\}$, i.e., fabricate a proof of the form $\{(u_i, s_i)\}_{i \in Q}$ that verifies. It does not make sense to consider adaptive adversaries because the proving phase consists of decoupled packets sent to the verifier; since no ephemeral quantities are exchanged between the shareholders, \mathcal{A} has no chance to collect potentially useful information and decide which of them to corrupt adaptively. We can thus confine ourselves to the static model, assuming that the attacker decides from the outset the set of corrupted parties.³ One could further assume that \mathcal{A} knows the public keys of individual shareholders; however, these are not advertised or used anyhow, although trivially inferred from protocol transcripts. Letting \mathcal{A} know them from the outset corresponds to a scenario where the public shares become advertised within the group for at least some initial timespan (e.g., for recognition and verification purposes) and \mathcal{A} compromises the device of a shareholder who has not yet erased them. Still, the security proof presented in Section 5.6 remains intact under this stronger assumption (cf. Remark TODO),

³Adaptive corruption during the execution of an equivalent DKG protocol in place of Shamir's sharing falls under the static model. TODO

indicating that an adversary who controls devices in the erasure-free model has no greater advantage than an adversary who only steals keys.

Let $(x_1, \dots, x_n; y)$ be the given (n, t) -sharing of some secret x and \mathcal{A} be an adversary who knows the shares $\{x_i\}_{i \in J}$ with $|J| = t - 1$. Let $Q \subset \{1, \dots, n\}$ be the shareholders that fall victim to impersonation under a purportedly valid proof $\{(u_i, s_i)\}_{i \in Q}$ fabricated by \mathcal{A} . Without loss of generality, $Q \subset J$ or $J \subset Q$. Indeed, since the black-box adversary knows nothing special about the uncorrupted shareholders, the choice of Q is irrelevant (except for its size) in the extent to which Q intersects with $\{1, \dots, n\} \setminus J$; the adversary can change Q without this affecting its strategy, provided that $|Q|$ is preserved and its original intersection with J remains included. We conclude with the following attack model.

- *Corruption and attack statement:*
 - \mathcal{A} chooses $J \subset \{1, \dots, n\}$ with $|J| = t - 1$.
 - \mathcal{A} chooses $Q \subset \{1, \dots, n\}$ with $Q \subset J$ or $J \subset Q$ and sends it along with J to its challenger \mathcal{C} .
 - \mathcal{C} runs $(x_1, \dots, x_n; y) \leftarrow D(x)$, $x \leftarrow_{\$} \mathbb{Z}_q$ and sends $y, \{x_i\}_{i \in J}$ to \mathcal{A} .
 - *Impersonation:* \mathcal{A} outputs a proof of the form $\sigma = \{(u_i, s_i)\}_{i \in Q}$.
- \mathcal{A} wins if σ is valid against y .

Successful impersonation is equivalent to condition (3.23) for the output σ .

ORST security against impersonation attacks should mean that every polynomial-time adversary wins the above game with at most negligible chances. The pragmatic requirement behind is that, if \mathcal{A} outputs a proof $\{(u_i, s_i)\}_{i \in Q}$ that verifies, then it should necessarily know x_j for some $j \notin J$ (or, equivalently, x). In other words, impersonation should be infeasible without knowledge of at least t secret shares. We will informally refer to this design requirement as the *knowledge property*.

This is instance of a more general pattern. Security notions for Schnorr-based protocols are naturally intermediated by some contextual knowledge property like above, which is the pragmatic security requirement per se. Consequently, a security reduction should –most probably– proceed by recovery of some contextual key in case of successful attack; by *knowing* this key, the adversary solves the DL or related hard problem, meaning that it cannot have succeeded (with non-negligible chances) in its attack. Key-recovery proceeds as follows. By rewinding the successful adversary at some special point and let it repeat successfully its attack, the rewinder *extracts* the key from the partially differing transcripts of the two successful attacks. Security reduces to the hardness assumption only because extraction yields the solution to the underlying hard problem under any circumstances that allow rewinding; if extraction is not possible under certain circumstances, security is –most probably– unprovable without further assumptions. It should be stressed that extractability is only a way to formalize the knowledge property in a security reduction argument; in particular, the knowledge property can hold true and pragmatically guarantee security without security being provable under the fixed hardness assumption.

Regarding ORST, the relevant extraction formula for x_j is evidently equation (3.37) from the key-recovery attack of Section 3.2.3. We will see (cf. Section 4.1.2) that this formula is insufficient under certain circumstances, indicating that ORST security is –most probably– unprovable under DL hardness alone. We would need

extra assumptions in order for security to become provable or, what is equivalent, reduce it to an easier problem than DL.

4.1.2. Extractability and $(t - 1)$ -OMDL hardness. Consider the following situation with $Q = J \cup \{j\}$, $|J| = t - 1$. The shareholders $i \in J$ engage in a session with a verifier and an adversary \mathcal{A} who knows $\{x_i\}_{i \in J}$ intercepts their communication; in particular, \mathcal{A} captures the transcript u_i, c_i, s_i , $i \in J$. Subsequently, \mathcal{A} fabricates a commitment u_j according to some-black-box strategy and sends it to the verifier, who responds with a challenge c_j . Finally, \mathcal{A} fabricates a response s_j according to some black-box strategy and sends it to the verifier. If the protocol is secure, the transcript u_i, c_i, s_i , $i \in Q$ should verify with at most negligible chances, even though \mathcal{A} knows x_i, u_i, c_i, s_i , $i \in J$ while fabricating u_j and s_j . Pragmatically, if the transcript verifies, then \mathcal{A} must necessarily know x_j .

Extraction is here concretely instantiated as follows. Let \mathcal{A}^* be an algorithm with rewindable black-box access to \mathcal{A} . If \mathcal{A} fabricates u_j, s_j such that the transcript u_i, c_i, s_i , $i \in Q$ verifies with probability ε , \mathcal{A}^* rewinds \mathcal{A} exactly before sending u_j and lets it complete a second impersonation attempt. By the forking lemma (cf. [12], or the reset lemma [9]), \mathcal{A} succeeds again with probability $\varepsilon^* \approx \varepsilon^2$. Let u_i^*, c_i^*, s_i^* , $i \in Q$ be the transcript of the second impersonation attack; since rewinding occurs at the last commitment, we have

$$(u_i^* = u_i \wedge c_i^* = c_i \wedge s_i^* = s_i) \forall i \in J \wedge u_j^* = u_j \wedge c_j^* \neq c_j \quad (4.1)$$

where the last condition holds with overwhelming probability true. Extractability means that \mathcal{A}^* *should* be able to recover x_j from x_i, u_i, c_i, s_i , $i \in J$ and u_j, c_j^*, s_j^* . We contend that this is impossible *without* further assumptions.

Condition (4.1) allows us to apply the computations of Section 3.2.3 in the present context, meaning that the relevant extraction formula is the same as (3.37). In particular, x_j is extractable in the form

$$x_j = \frac{s_j - s_j^*}{c_j - c_j^*} + \frac{1}{\mu_j} \sum_{i \in J} (\mu_i - \mu_i^*) (s_i - r_i - c_i x_i), \quad r_i \equiv \log u_i, \quad (4.2)$$

where we have further applied the fact $s_i^* = s_i, \forall i \in J$. However, \mathcal{A}^* cannot know the logarithms r_i , $i \in J$ of the intercepted commitments and consequently cannot use this formula to recover x_j . Extraction is not possible in the present context.

Since we have found circumstances where extraction cannot complete without further assumptions, security is –most probably– unprovable without further assumptions as well. In particular, ORST security against impersonation attacks is –most probably– *not* reducible to discrete-logarithm (DL) hardness alone. However, \mathcal{A}^* would indeed be able to extract x_j if it could somehow see the logarithms of the intercepted commitments, e.g., if it controlled the respective shareholders' devices, or could at least subvert their random number generator or other security sensitive module. We can codify such scenarios under the assumption that \mathcal{A}^* is allowed to issue at most $t - 1$ queries to a hypothetical DL oracle; extraction of x_j by means of (4.2) would then become automatically possible, suggesting that security against impersonation attacks might as well be reducible to the $(t - 1)$ -OMDL hardness assumption (cf. Section 4.3). Note that $(t - 1)$ -OMDL is indeed an easier problem than DL, plausibly yielding the extra assumptions needed for security provability.

4.1.3. *Implications of $(t - 1)$ -OMDL hardness.* The black-box pattern of the previous section is rather restricted as compared to the general attack model for ORST (cf. Section 4.1.1). This is because \mathcal{A} does not influence the compromised shareholders, mounting its actual attack only after collecting their transcripts; in particular, the fabrication of u_j and s_j is deferred until after $u_i, s_i, i \in J$ have resolved, in order for extractability failure to be demonstrated in the simplest possible way (i.e., rewinding instead of some more complicated forking). At the same moment, the unprovability argument indicates in which way security becomes provable under the stronger $(t - 1)$ -OMDL hardness assumption. Specifically, in order to be able to apply extraction, it seems that \mathcal{A}^* should at least reproduce the conditions of Remark 3.2.3 and let \mathcal{A} operate in that context. While this remains more restricted than the general attack model of Section 4.1.1, it generalizes by far the restricted attack pattern of Section 4.1.2, allowing \mathcal{A} to fabricate u_j before the rest shareholders send their commitments and moreover control their responses. We propose the following preliminary attack model for \mathcal{A} , representing a tradeoff between the general attack model and security provability under the $(t - 1)$ -OMDL hardness assumption.

- *Corruption and attack statement:*
 - \mathcal{A} chooses $J \subset \{1, \dots, n\}$ with $|J| = t - 1$ and sends it along with some $Q \subset \{1, \dots, n\}$ to its challenger \mathcal{C} .
 - \mathcal{A} chooses $Q \subset \{1, \dots, n\}$ with $Q \subset J$ or $J \subset Q$ and sends it along with J to its challenger \mathcal{C} .
 - \mathcal{C} runs $(x_1, \dots, x_n; y) \leftarrow D(x)$, $x \leftarrow_{\mathbb{S}} \mathbb{Z}_q$ and sends $y, \{x_i\}_{i \in J}$ to \mathcal{A} .
- *Interception:* \mathcal{C} samples $u_i \leftarrow_{\mathbb{S}} \mathbb{G}$, $i \in J$ and sends $\{u_i\}_{i \in J}$ to \mathcal{A} .
- *Impersonation:* \mathcal{A} outputs a proof of the form $\{(u_i, s_i)\}_{i \in Q}$.
 \mathcal{A} wins if $\{(u_i, s_i)\}_{i \in Q}$ verifies against y .

ORST is meant to be secure against impersonation attacks if every polynomial-time adversary wins this game with at most negligible advantage.

Observe that no restrictions are yet imposed on $|Q|$, allowing the possibility that \mathcal{A} fabricates a valid proof with less than t components. This is plausible because proofs of this kind are known to exist in abundance (cf. Remarks 3.29 and 3.32). However, under the $(t - 1)$ -OMDL hardness assumption, fabricating such proofs can be ruled out as infeasible. Let \mathcal{A} be a polynomial-time adversary that wins with $Q \subset J$ (i.e., $|Q| < t$). Let $J = \{i_1, \dots, i_{t-1}\}$ with $i_1 < \dots < i_{t-1}$, so that $Q = \{i_1, \dots, i_m\}$ for some $m \in \{1, \dots, t - 1\}$. We will construct a $(t - 1)$ -OMDL adversary \mathcal{A}^* as a wrapper of \mathcal{A} with equal time complexity and advantage. When given w_0, \dots, w_{t-1} by its challenger (cf. Game 4.7), \mathcal{A}^* sets $y \leftarrow w_0$ and $u_{i_k} \leftarrow w_k$, $1 \leq k \leq t - 1$. Next, \mathcal{A}^* samples $x_i \leftarrow_{\mathbb{S}} \mathbb{Z}_q$, $i \in J$ and forwards $y, \{x_i\}_{i \in J}$ to \mathcal{A} . By the security property of Shamir's sharing (cf. Remark A.5), y and $\{x_i\}_{i \in J}$ determine some (n, t) -sharing x_1, \dots, x_n of $x \equiv \log y$ indistinguishably, meaning that \mathcal{A}^* simulates perfectly the corruption phase of the game. Subsequently, \mathcal{A}^* sends $\{u_i\}_{i \in Q}$ to \mathcal{A} and lets it proceed to impersonation; since its challenger has been perfectly simulated, \mathcal{A} outputs a valid proof $\{(u_i, s_i)\}_{i \in Q}$ with probability equal to its advantage. Condition (3.23) is then satisfied and \mathcal{A}^* can recover x by issuing at most $t - 1$ discrete-logarithm queries. More accurately, \mathcal{A}^* queries

$$r_{i_k} \leftarrow \mathcal{O}_{\mathbb{G}}^{\text{dlog}}(u_{i_k}), \quad 1 \leq k \leq t - 1$$

and uses r_{i_k} , $1 \leq k \leq m$ to compute x . Finally, \mathcal{A}^* sets $z_0 \leftarrow x$, and $z_k \leftarrow r_{i_k}$ for $1 \leq k \leq t-1$, and wins by forwarding z_0, \dots, z_{t-1} to its challenger. In short, if \mathcal{A} succeeds with non-negligible probability, then \mathcal{A}^* solves the $(t-1)$ -OMDL problem with equal advantage, which contradicts the hardness assumption.

This straight-line reduction shows that we can restrict attention to the case $J \subsetneq Q$ (i.e. $|Q| \geq t$), excluding the possibility of impersonating less than t shareholders (without knowledge of at least t secret shares). This is evidently a security requirement for the ORST protocol, captured (as should) by the $(t-1)$ -OMDL hardness assumption. On the other hand, TODO

4.2. Security against impersonation attacks. TODO We conclude with the following model for impersonation attacks against the ORST protocol. Given a cyclic group $\mathbb{G} = \langle g \rangle$ of order q and integers $1 \leq t \leq n < q$, we define the impersonation attack against ORST to be the following interactive game.

Game 4.3. ($\text{ORST}_{\mathbb{G},n,t}^{\text{imp}} - \text{Impersonation}$) Given adversary \mathcal{A} and challenger \mathcal{C} ,

- *Corruption and attack statement:*
 - \mathcal{A} chooses $J \subset \{1, \dots, n\}$ with $|J| = t-1$ and sends it along with some $j \in \{1, \dots, n\} \setminus J$ to the challenger. Denote $Q = J \cup \{j\}$.
 - \mathcal{C} runs $(x_1, \dots, x_n; y) \leftarrow D(x)$, $x \leftarrow_{\mathbb{S}} \mathbb{Z}_q$ and sends $y, \{x_i\}_{i \in J}$ to \mathcal{A} .
- *Interception:* \mathcal{C} samples $u_i \leftarrow_{\mathbb{S}} \mathbb{G}$, $i \in J$ and sends $\{u_i\}_{i \in J}$ to \mathcal{A} .
- *Impersonation:* \mathcal{A} outputs a proof of the form $\sigma = \{(u_i, s_i)\}_{i \in Q}$.

\mathcal{A} wins if $\mathcal{V}(y, \sigma) = \text{true}$.

The probability that \mathcal{A} wins Game 4.3 is denoted by

$$\text{Adv}_{\mathbb{G},n,t}^{\text{ORST}}[\mathcal{A}]. \quad (4.4)$$

ORST *security in \mathbb{G}* is the assumption that (4.4) is negligible for every choice of the parameters n, t and every polynomial-time adversary \mathcal{A} .

Remark 4.5. (Adversarial components) By the security property of Shamir's secret sharing (cf. Remark A.5), the parameters j and J (and consequently Q) may be regarded as fixed for \mathcal{A} . In particular, we can regard \mathcal{A} as a pair of non-interactive algorithms $(\mathcal{A}_0, \bar{\mathcal{A}})$ of the form

$$(j, J) \leftarrow \mathcal{A}_0(\cdot), \quad \{(u_i, s_i)\}_{i \in Q} \leftarrow \bar{\mathcal{A}}(y, \{x_i\}_{i \in J}, \{u_i\}_{i \in J})$$

with the latter consuming the adversarial view during the impersonation phase in order to fabricate a commitment u_j and responses $\{s_i\}_{i \in Q}$ such that the output $\{(u_i, s_i)\}_{i \in Q}$ verifies. Note that, when using \mathcal{A} in a black-box reduction, the logarithm of the fabricated commitment u_j cannot be assumed to be known and we will normally have to eliminate it from our computations.

Remark 4.6. (Deterministic adversaries) For every adversary \mathcal{A} attacking Game 4.3, there exists a deterministic adversary \mathcal{A}' of equal time complexity such that

$$\text{Adv}_{\mathbb{G},n,t}^{\text{ORST}}[\mathcal{A}'] \geq \text{Adv}_{\mathbb{G},n,t}^{\text{ORST}}[\mathcal{A}].$$

This follows from a generic argument regarding any sufficiently simple adversary \mathcal{A} attacking a game G . Consider the residual deterministic strategies $\mathcal{A}_1, \dots, \mathcal{A}_m$ obtained by fixing the possible values of \mathcal{A} 's random coins and let p_i be the probability that \mathcal{A} runs \mathcal{A}_i , which is itself an adversary for G . We have

$$\begin{aligned} \text{Adv}^G[\mathcal{A}] &= \sum_{i=1}^m \Pr[\mathcal{A}_i \text{ wins } G \wedge \mathcal{A} \text{ runs } \mathcal{A}_i] \\ &= \sum_{i=1}^m \Pr[\mathcal{A} \text{ runs } \mathcal{A}_i] \cdot \Pr[\mathcal{A}_i \text{ wins } G \mid \mathcal{A} \text{ runs } \mathcal{A}_i] \\ &= \sum_{i=1}^m p_i \cdot \text{Adv}^G[\mathcal{A}_i] \end{aligned}$$

Since $p_1 + \dots + p_m = 1$, we get $\text{Adv}^G[\mathcal{A}_i] \geq \text{Adv}^G[\mathcal{A}]$ for some $i \in \{1, \dots, m\}$.

4.3. One-more discrete-logarithm (OMDL) hardness. The OMDL problem family is a parametrized extension of the discrete-logarithm (DL) problem, yielding a sequence of increasingly easier related problems. It was introduced in [9], [10] as a way to capture protocol design properties that pragmatically guarantee security under DL hardness while most probably being themselves irreducible to the latter. Note that reducing security to an easier problem, (i.e., a stronger hardness assumption) amounts usually to assumptions regarding the black-box adversary (i.e., a narrower threat model). This represents a tradeoff between the high plausibility of the weaker assumption and having a concrete security proof at hand, which may otherwise be impossible. Specifically, under the extra assumptions a formal security proof may become possible, while a proof under the weaker assumption may not even exist; since the solvability of the harder problem implies solvability of the easier one, the restricted security proof may provide good evidence that security holds true under the weaker assumption.
 TODO The tradeoff is often made in the context of secure multi-party protocols, where the adversary has extended control on its execution environment through the corrupted shareholders.
 TODO ref to our case

The *discrete-logarithm* (DL) *oracle* for $\mathbb{G} = \langle g \rangle$ is a hypothetical polynomial-time procedure of the form $z \leftarrow \mathcal{O}_{\mathbb{G}}^{\text{dlog}}(u)$, $u \in \mathbb{G}$, such that $u = g^z$. Given $\mathcal{O}_{\mathbb{G}}^{\text{dlog}}$ and integer $t \geq 1$, we consider the following problem.

Game 4.7. (OMDL $_{\mathbb{G}, t-1}$ – *One-more discrete-logarithm problem*) Given an adversary \mathcal{A} with challenger \mathcal{C} ,

- \mathcal{C} samples $w_i \leftarrow_{\$} \mathbb{G}$, $0 \leq i \leq t-1$ and sends them to \mathcal{A}
- \mathcal{A} outputs z_0, \dots, z_{t-1}

\mathcal{A} wins if $\bigwedge_{i=0}^{t-1} w_i = g^{z_i}$ and it has issued at most $t-1$ queries to $\mathcal{O}_{\mathbb{G}}^{\text{dlog}}$.

The probability that \mathcal{A} solves OMDL $_{\mathbb{G}, t-1}$ is denoted by

$$\text{Adv}_{\mathbb{G}, t-1}^{\text{OMDL}}[\mathcal{A}]. \quad (4.8)$$

$(t-1)$ -OMDL *hardness in* \mathbb{G} is the assumption that (4.8) is negligible for every polynomial-time adversary. Evidently, 0-OMDL is equivalent to DL. If an adversary

solves $\text{ORST}_{\mathbb{G}, t-1}$ for $t < t^*$, then it trivially solves $\text{ORST}_{\mathbb{G}, t^*-1}$, i.e., (t^*-1) -ODML is generally an easier problem than $(t-1)$ -OMDL. In particular, $(t-1)$ -OMDL, $t > 1$ is generally a stronger assumption than DL.

5. SECURITY PROOF

We prove ORST security against impersonation attacks (cf. Section 4.2) under OMDL hardness (cf. Section 4.3) in the random oracle model. The rationale behind the attack model is found in Section 4.1.1.

Our reduction generalizes the rewinding argument of the non-distributed case in the threshold setting in order to extract the unknown share x_j . This is a non-trivial generalization. As indicated by the key-recovery attack of Section 3.2.3, the extractor comes with a residue stemming from the $t-1$ degrees of freedom in perturbing a valid proof with fixed commitments (cf. Remark 3.29); it becomes useful only if the logarithms of the intercepted commitments are known to the reduction, which is the reason for reducing to $(t-1)$ -OMDL hardness (cf. Section 4.1.2). Moreover, extraction presupposes the fixture of all but one challenges, making rewinding useless in the threshold setting. We overcome this obstacle by employing the Local Forking Lemma (cf. Section B) and ensuring that the conditions for its applicability are satisfied. The attack simulator is presented in Section 5.4 and its components are developed in Sections 5.1 to 5.3.

5.1. Index mapping. Given $J \subset \{1, \dots, n\}$, we can represent any collection $\alpha_1, \dots, \alpha_m$ with $|J| = m$ in the form $\{\beta_i\}_{i \in J}$ using the Zip operator below.

Algorithm 1 $\text{Zip}(\alpha_1, \dots, \alpha_m; J)$	Algorithm 2 $\text{Zip}^{-1}(\{\beta_i\}_{i \in J})$
$\{i_1, \dots, i_m\} \leftarrow J$ with $i_1 < \dots < i_m$	$\{i_1, \dots, i_m\} \leftarrow J$ with $i_1 < \dots < i_m$
for $k = 1, \dots, m$	for $k = 1, \dots, m$
$\beta_{i_k} \leftarrow \alpha_k$	$\alpha_k \leftarrow \beta_{i_k}$
return $\{\beta_i\}_{i \in J}$	return $\alpha_1, \dots, \alpha_m$

It should be clear that Zip preserves the uniform sampling.

Lemma 5.1. Given $J \subset \{1, \dots, n\}$, $|J| = m \geq 1$, the probability distribution

$$\{\text{Zip}(\alpha_1, \dots, \alpha_m; J) : \alpha_i \leftarrow_{\S} A, i = 1, \dots, m\}$$

is identical to the distribution $\{\{\beta_i\}_{i \in J} : \beta_i \leftarrow_{\S} A, i \in J\}$.

Proof. Follows directly from the definition of Zip. \square

5.2. Shamir sharing simulation. Let $y^* \in \mathbb{G}$ and $x_1^*, \dots, x_{t-1}^* \in \mathbb{Z}_q$, $t \leq n$. Given $J \subset \{1, \dots, n\}$ with $|J| = t-1$ and $j \in \{1, \dots, n\} \setminus J$, we define

$$(\{x_i\}_{i \in J}, y) \leftarrow \mathcal{D}(y^*, x_1^*, \dots, x_{t-1}^*; j, J) \quad (5.2)$$

to be Algorithm 3.

Algorithm 3 $\mathcal{D}(y^*, x_1^*, \dots, x_{t-1}^*; j, J)$

```

1: Set  $y_j \leftarrow y^*$ 
2:  $\{x_i\}_{i \in J} \leftarrow \text{Zip}(x_1^*, \dots, x_{t-1}^*)$ 
3: Set  $y_i \leftarrow g^{x_i}$ ,  $i \in J$ 
4:  $Q \leftarrow J \cup \{j\}$ 
5:  $y \leftarrow \prod_{i \in Q} y_i^{\lambda_i}$ 
6: return  $\{x_i\}_{i \in J}, y$ 

```

Its properties are summarized as follows.

Lemma 5.3. *Let $J \subset \{1, \dots, n\}$, $|J| = t - 1$ with $t \leq n$, and $j \in \{1, \dots, n\} \setminus J$. Given (5.2), there exists a unique (n, t) -sharing $(x_1, \dots, x_n; y)$ such that*

$$\{x_i\}_{i \in J} = \text{Zip}(x_1^*, \dots, x_{t-1}^*) \wedge y^* = g^{x_j} \quad (5.4)$$

Moreover, the probability distribution

$$\{(\{x_i\}_{i \in J}, y) : \left\{ \begin{array}{l} (\{x_i\}_{i \in J}; y) \leftarrow \mathcal{D}(y^*, x_1^*, \dots, x_{t-1}^*; j, J) \\ y^* \leftarrow_{\$} \mathbb{G}, \quad x_i^* \leftarrow_{\$} \mathbb{Z}_q, \quad 1 \leq i \leq t-1 \end{array} \right\}\}$$

is identical to $\{(\{x_i\}_{i \in J}, y) : (x_1, \dots, x_n; y) \leftarrow D(x), x \leftarrow_{\$} \mathbb{Z}_q\}$.

Proof. By the elementary properties of interpolation, x_1^*, \dots, x_{t-1}^* and y^* uniquely determine a sharing subject to the constraint (5.4), with the outputs of \mathcal{D} mapping to those of D bijectively. Equality of distributions follows from the security of the Shamir's sharing (cf. Remark A.5). \square

5.3. Hash oracles. When querying a random oracle H (cf. Algorithm 5) in a simulation context, we usually need to keep track of distinct oracle queries. Given a finite tape of preallocated hashes $h = [h_1 \dots, h_m]$ and initially empty lookup tables $\text{Map}_1, \text{Map}_2$, we define the *hash oracle* $\mathcal{O}^{\text{hash}}$ to be Algorithm 4.

Algorithm 4 $\mathcal{O}^{\text{hash}}(u)$

```

if  $u \notin \text{Dom}(\text{Map}_1)$ 
   $c \leftarrow h[\nu]$ 
   $\text{Map}_1[u] \leftarrow c$ 
   $\text{Map}_2[u] \leftarrow \nu$ 
   $\nu \leftarrow \nu + 1$ 
else
   $c \leftarrow \text{Map}_1[u]$ 
return  $c$ 

```

Algorithm 5 $H(u)$

```

if  $u \notin \text{Dom}(\text{Map})$ 
   $c \leftarrow_{\$} C$ 
   $\text{Map}[u] \leftarrow c$ 
else
   $c \leftarrow \text{Map}[u]$ 
return  $c$ 

```

That is, $\mathcal{O}^{\text{hash}}$ operates like a random oracle with lookup table Map_1 but draws hashes from h instead of sampling them uniformly; upon inserting a fresh hash to Map_1 , it takes care to (1) update the tape pointer so that the next predefined hash be available on demand and (2) insert the current pointer value to Map_2 so that the respective query can be easily tracked by order. Evidently, if h_1, \dots, h_n are uniformly sampled from the codomain of H , then $\mathcal{O}^{\text{hash}}$ cannot be distinguished from H . We state this remark formally for the sake of reference.

Lemma 5.5. (*Random oracle simulation*) *Given a random oracle H with codomain C and empty lookup table Map , initialize a hash oracle $\mathcal{O}^{\text{hash}}$ over an empty tape h of size $m \geq 1$ and initially empty lookup tables Map_1 and Map_2 . For any distinct elements u_1, \dots, u_m , the probability distribution*

$$\{(c_1, \dots, c_m) : c_\nu \leftarrow \mathcal{O}^{\text{hash}}(u_\nu), h[\nu] \leftarrow_{\$} C, \nu = 1, \dots, m\}$$

is identical to the distribution $\{(c_1, \dots, c_m) : c_\nu \leftarrow H(u_\nu), \nu = 1, \dots, m\}$.

Proof. Follows directly from the definition of $\mathcal{O}^{\text{hash}}$. □

5.4. Simulated impersonation attack. We define an algorithm $\mathcal{S}[\mathcal{A}]$ that simulates Game 4.3 with adversary \mathcal{A} if H is a random oracle. Roughly speaking, it accepts as input the ingredients for the adversarial view and outputs the fabricated proof along with a success index for \mathcal{A} 's impersonation attempt. The simulation is perfect meaning that, if run with uniform input, the probability of success equals the adversary's advantage of winning the real game.

More accurately, given a black-box adversary \mathcal{A} to game $\text{ORST}_{\mathbb{G}, n, t}$, the desired simulation is wrapper of \mathcal{A} of the form

$$(k, (j, \sigma)) \leftarrow \mathcal{S}[\mathcal{A}](y^*, x_1^*, \dots, x_{t-1}^*, w_1, \dots, w_{t-1}; h_1, \dots, h_m),$$

arranging its input values appropriately in order to emulate \mathcal{A} 's execution environment. The values $y^*, x_1^*, \dots, x_{t-1}^*$ will uniquely determine the attacked (n, t) -sharing, yielding the corrupted shares in the process. Similarly, w_1, \dots, w_{t-1} will be used as the intercepted commitments. Next, the embedded replica of \mathcal{A} consumes its view in order to fabricate a purportedly valid proof σ with the difference that, instead of invoking H directly, it queries an embedded hash oracle $\mathcal{O}^{\text{hash}}$ with pre-allocated hashes h_1, \dots, h_m (cf. Section 5.3). If needed, \mathcal{A} may be reprogrammed without loss of generality to query hashes for all the commitments it sees (including the one fabricated by itself); in this case $m \geq 1$. Finally, an ORST -verifier \mathcal{V} is included with access to the same hash oracle for verifying the output of the embedded adversary. The exact procedure is Algorithm 6.

Algorithm 6 $\mathcal{S}[\mathcal{A}](y^*, x_1^*, \dots, x_{t-1}^*, w_1, \dots, w_{t-1}; h_1, \dots, h_m)$

```

1:  $h \leftarrow [h_1, \dots, h_m]$ ,  $\nu \leftarrow 1$ ,  $\text{Map}_1 = \emptyset$ ,  $\text{Map}_2 = \emptyset$  // Setup  $\mathcal{O}^{\text{hash}}$ 
2:  $(j, J) \leftarrow \mathcal{A}_0(\cdot)$  // Attack statement
3:  $(\{x_i\}_{i \in J}, y) \leftarrow \mathcal{D}(y^*, x_1^*, \dots, x_{t-1}^*; j, J)$  // Corruption
4:  $\{u_i\}_{i \in J} \leftarrow \text{Zip}(w_1, \dots, w_{t-1})$  // Interception
5:  $\{(u_i, s_i)\}_{i \in Q} \leftarrow \bar{\mathcal{A}}(y, \{x_i\}_{i \in J}, \{u_i\}_{i \in J})$  // Impersonation
6:  $\text{res} \leftarrow \mathcal{V}(y, \{(u_i, s_i)\}_{i \in Q})$ 
7: for  $i \in Q$ 
8:    $c_i \leftarrow \text{Map}_1[u_i]$ 
9:  $\sigma \leftarrow \{(u_i, c_i, s_i)\}_{i \in Q}$ 
10:  $k \leftarrow \text{Map}_2[u_j]$ 
11: if  $\text{res} = \text{true}$ 
12:   return  $(k, (j, \sigma))$ 
13: return  $(0, (j, \sigma))$ 

```

$\mathcal{O}^{\text{hash}}$ is first initialized with tape $[h_1, \dots, h_m]$ and its lookup tables set to empty. After \mathcal{A} makes its attack statement, the virtual challenger accordingly simulates the unique sharing $(x_1, \dots, x_n; y) \leftarrow D(x)$, $x \leftarrow \mathbb{Z}_q$ subject to the constraint (5.4); in particular, $x_j = \log y^*$. It then generates commitments $\{u_i\}_{i \in J}$ and feeds them along with y and $\{x_i\}_{i \in J}$ to \mathcal{A} , who proceeds to impersonation and generates a proof of the expected form; in particular, $Q = J \cup \{j\}$. Note that all oracle queries are implicitly issued at step 5, including the challenges $c_i \leftarrow \mathcal{O}^{\text{hash}}(u_i)$, $i \in Q$, which are collected and attached to the fabricated proof for later usage (steps 7-9). We also locate the unique $k \in \{1, \dots, m\}$ for which $c_j = h[k]$ (step 10), indicating that the fabricated commitment was submitted at the k -th distinct query. If \mathcal{A} succeeds in its impersonation attempt, the simulation returns $(k, (j, \sigma))$; otherwise, it returns $(0, (j, \sigma))$ to indicate failure.

Proposition 5.6. (ORST $_{\mathbb{G}, n, t}$ – Perfect simulation) *If H is modelled as a random oracle, $\mathcal{S}[\mathcal{A}]$ simulates Game 4.3 perfectly if invoked with uniformly random input, in the sense that the real and simulated game executions are indistinguishable from the adversary’s viewpoint. More accurately, the probability distribution of \mathcal{A} ’s view in the real attack game is identical to the distribution of \mathcal{A} ’s input in*

$$\begin{aligned}
(k, (j, \sigma)) &\leftarrow \mathcal{S}[\mathcal{A}](y^*, x_1^*, \dots, x_{t-1}^*, w_1, \dots, w_{t-1}; h_1, \dots, h_m), \\
y^* &\leftarrow_{\$} \mathbb{G}, \quad x_i^* \leftarrow_{\$} \mathbb{Z}_q, \quad w_i \leftarrow_{\$} \mathbb{G}, \quad 1 \leq i \leq t-1, \\
h_\nu &\leftarrow_{\$} \mathbb{Z}_q, \quad 1 \leq \nu \leq m
\end{aligned} \tag{5.7}$$

In the context of this simulation, if $\bar{\mathcal{A}}$ succeeds in its impersonation attempt, i.e., $k \geq 1$, the input value h_k maps to $c_j \equiv H(u_j)$ of the real game execution.

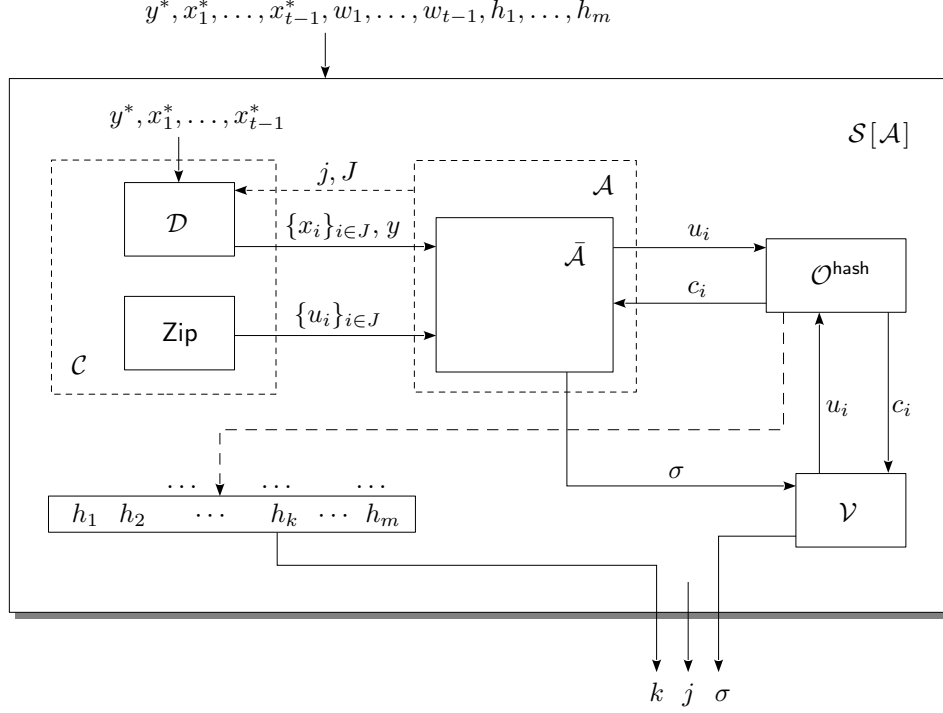


FIGURE 1. The $\text{ORST}_{\mathbb{G},n,t}$ simulator

Proof. In the real game execution, the adversarial view consists of the corrupted shares $\{x_i\}_{i \in J}$ along with the combined public key y , the intercepted commitments $\{u_i\}_{i \in Q}$ and the hashes queried during the impersonation phase. These three blocks are independent from each other; clearly, the first and second blocks are independent, while the third block is independent from the others because H is a random oracle. It thus suffices to check that the respective blocks have identical distributions in the simulated execution (5.7). The first block is perfectly simulated due to Lemma 5.3. The second block is also perfectly simulated due to Lemma 5.1. Finally, the third block is perfectly simulated due to Lemma 5.5. The fact that $h_k = h[k]$ corresponds to $H(u_j)$ follows directly from step 10, Algorithm 6. \square

We denote by $\text{acc}(\mathcal{S}[\mathcal{A}])$ the probability that the embedded adversary succeeds in the context of a simulated attack when run with uniformly random input (cf. (5.7)). This is related to the real adversarial advantage as follows.

Corollary 5.8. *Given an adversary \mathcal{A} for Game 4.3,*

$$\text{acc}(\mathcal{S}[\mathcal{A}]) = \text{Adv}_{\mathbb{G},n,t}^{\text{ORST}}[\mathcal{A}] \quad (5.9)$$

Proof. Direct consequence of Proposition 5.6. \square

5.5. Local fork and extraction. In the non-distributed case $t = 1$, the simulator $\mathcal{S}[\mathcal{A}]$ yields the execution environment of the rewindable adversary \mathcal{A} that is usually

employed to reduce the security of ordinary Schnorr identification to DL hardness in the random oracle model. In the non-distributed case, the extractor is generalized as (3.37) provided that the forked attack runs with the same hashes except for $c_j \leftarrow H(u_j)$. This implies that the rewinding argument cannot be effective in the threshold setting; indeed, resetting the attack before \mathcal{A} sends u_j does not guarantee that subsequent oracle responses remain the same and, in fact, they differ almost certainly. Nevertheless, when invoking $\mathcal{S}[\mathcal{A}]$ we predefine the hashes h_1, \dots, h_m returned by the embedded oracle in respective order, which allows to rerun the simulator with the same hashes except for $c_j = h_k$. In this case, the Local Forking Lemma (cf. Lemma B.5) becomes applicable in place of the general forking lemma (cf. [12] or reset lemma, [9]), provided that \mathcal{A} is deterministic.

Recall that the local fork of $\mathcal{S}[\mathcal{A}]$ is an algorithm of the form

$$(b, \text{out}) \leftarrow \text{LocalFork}_{\mathcal{S}[\mathcal{A}]}(y^*, x_1^*, \dots, x_{t-1}^*, w_1, \dots, w_{t-1})$$

(cf. Appendix B) operating as follows. If the embedded adversary \mathcal{A} succeeds in its first impersonation attempt, the fork reexecutes the simulation with the same input values except for the k -th predefined hash, which is afresh uniformly sampled as h_k^* ; otherwise, the fork aborts ($b = 0$). When the second simulation completes, the fork checks if \mathcal{A} succeeds again under the condition $h_k^* \neq h_k$; in this case, it returns the successive outputs $\tau \equiv (j, \sigma)$ and $\tau^* \equiv (j, \sigma^*)$ of both simulations; otherwise, it aborts ($b = 0$).

Evidently, $\text{forkAcc}(\mathcal{S}[\mathcal{A}])$ (cf. def. (B.4)) expresses here the probability that \mathcal{A} succeeds both in its first and forked impersonation attempt when consuming the same uniformly random view (except for the k -th distinct hash). In this case, the outputs of the two simulated attacks are correlated as follows.

Proposition 5.10. (*Local fork and extraction*) *Let \mathcal{A} be a deterministic adversary of Game 4.3 and $\tau = (j, \sigma)$, $\tau^* = (j, \sigma^*)$ be outputs of a non-aborting execution*

$$(1, (\tau, \tau^*)) \leftarrow \text{LocalFork}_{\mathcal{S}[\mathcal{A}]}(y^*, x_1^*, \dots, x_{t-1}^*, w_1, \dots, w_{t-1})$$

Then σ, σ^ are of the form*

$$\{(u_i, c_i, s_i)\}_{i \in Q}, \{(u_i, c_i^*, s_i^*)\}_{i \in Q} \quad (5.11)$$

respectively and subject to the constraint

$$c_j \neq c_j^* \wedge c_i = c_i^*, \forall i \in Q \setminus \{j\} \quad (5.12)$$

Moreover, the extraction formula

$$x_j = \frac{s_j - s_j^*}{c_j - c_j^*} + \frac{1}{\mu_j} \sum_{i \in J} \left(\mu_i (s_i - r_i - c_i x_i) - \mu_i^* (s_i^* - r_i - c_i x_i) \right) \quad (5.13)$$

for $x_j \equiv \log y^$ holds true, where $r_i \equiv \log u_i$ and $\{\mu_i\}_{i \in Q}$, $\{\mu_i^*\}_{i \in Q}$ are the weights of σ and σ^* respectively.*

Proof. Let $\sigma = \{(u_i, c_i, s_i)\}_{i \in Q}$, $\sigma^* = \{u_i^*, c_i^*, s_i^*\}_{i \in Q}$ so that $u_i^* = u_i, \forall i \neq j$. Fabricating the j -th commitment completes before submitting it to the hash oracle and is thus unaffected by the k -th hash. Since \mathcal{A} is deterministic and its view is the

same in both simulations (apart from the k -th hash), we conclude that $u_j = u_j^*$. The constraint regarding challenges follows because $c_j^* = h_k^* \neq h_k = c_j$ and the rest are drawn from $\{h_\nu\}_{\nu \neq k}$. We proceed to the extraction formula. Since $c_i = c_i^*$ for $i \neq j$, by definition of weights we get

$$\mu_j = \mu_j^* \quad (5.14)$$

(cf. Definition 3.8). Since both σ, σ^* are valid, Corollary 3.33 implies

$$\sum_{i \in Q} \mu_i (s_i - r_i - c_i x_i) = 0, \quad \sum_{i \in Q} \mu_i^* (s_i^* - r_i - c_i^* x_i) = 0.$$

Subtracting terms and applying (5.14) yields

$$\mu_j (s_j - s_j^* - (c_j - c_j^*) x_j) + \sum_{i \in J} (\mu_i (s_i - r_i - c_i x_i) - \mu_i^* (s_i^* - r_i - c_i^* x_i)) = 0.$$

Extraction follows because $c_j - c_j^* \neq 0$. \square

5.6. Main theorem. TODO: Refer to the appendix when needed: deterministic operators without narrowing the threat model

Theorem 5.15. (ORST – Security against impersonation attacks) $\text{ORST}_{\mathbb{G}, n, t}$ is secure against impersonation attacks under $(t - 1)$ -OMDL hardness in the random oracle model. In particular, if H is modelled as a random oracle, for every polynomial-time adversary \mathcal{A} that attacks $\text{ORST}_{\mathbb{G}, n, t}^{\text{imp}}$ there exists a polynomial-time adversary \mathcal{A}^* that solves $\text{OMDL}_{\mathbb{G}, t-1}$ with advantage

$$\text{Adv}_{\mathbb{G}, t-1}^{\text{OMDL}}[\mathcal{A}^*] \geq \frac{1}{q_{\text{ro}} + 1} \text{Adv}_{\mathbb{G}, n, t}^{\text{ORST}}[\mathcal{A}]^2, \quad (5.16)$$

where $q_{\text{ro}} + 1$ is the number of distinct queries issued by \mathcal{A} to the random oracle.

Proof. We construct \mathcal{A}^* as a wrapper of $\text{LocalFork}_{\mathcal{S}[\mathcal{A}]}$; after forking a successful impersonation attack, the wrapper should namely be able to recover the unknown share by issuing at most $t - 1$ discrete-logarithm queries and then use it to solve $\text{OMDL}_{\mathbb{G}, t-1}$ with comparable chances. We assume without loss of generality that \mathcal{A} is deterministic (cf. TODO).

When given w_0, w_1, \dots, w_{t-1} by its challenger (cf. Game 4.7), \mathcal{A}^* operates as shown in Algorithm 7. Observe that \mathcal{A}^* aborts if the fork aborts, that is, unless both attacks by \mathcal{A} are successful; otherwise, the fork fulfills the conditions of Proposition 5.10 and the proofs fabricated by \mathcal{A} are as shown at steps 7 and 8. By definition of $\mathcal{S}[\mathcal{A}]$ (cf. Algorithm 6), the implicit attacks of step 3 unfold against the sharing $(x_1, \dots, x_n; y)$ that is uniquely determined by the constraint

$$\{x_i\}_{i \in J} = \text{Zip}(x_1^*, \dots, x_{t-1}^*) \wedge y^* = g^{x_j} \quad (5.17)$$

(cf. Lemma 5.3). Similarly, the intercepted commitments u_i , $i \in J$ comprise of the input values w_1, \dots, w_{t-1} ; in particular,

$$\{u_i\}_{i \in J} = \text{Zip}(w_1, \dots, w_{t-1})$$

(cf. Algorithm 6, step 4). In steps 9-14 \mathcal{A}^* computes the parameters needed for the extraction at step 15. The compromised shares are recomputed at step 13 using

the first part of constraint (5.17), while the logarithms of u_i , $i \in J$ are retrieved at step 14 by submitting them to the DL oracle. Since $y^* = w_0$, Proposition 5.10 ensures that the computation of step 15 yields

$$z_0 = \log w_0 \tag{5.18}$$

Similarly, due to (5.6), the operation of step 17 yields

$$z_i = \log w_i, \quad 1 \leq i \leq t-1$$

In short, z_0, z_1, \dots, z_{t-1} is the solution to the original problem instance and \mathcal{A}^* has issued at most $t-1$ distinct queries to the DL oracle.

This discussion shows that \mathcal{A}^* solves $\text{OMDL}_{\mathbb{G}, t-1}$ exactly if it does not abort. Specifically, since w_0, w_1, \dots, w_{t-1} are uniformly sampled by its challenger, \mathcal{A}^* 's advantage is equal to the probability that

$$(b, \text{out}) \leftarrow \text{LocalFork}_{\mathcal{S}[\mathcal{A}]}(y^*, x_1^*, \dots, x_{t-1}^*, w_1, \dots, w_{t-1})$$

is a non-aborting execution when run with uniform input; or, in other words,

$$\text{Adv}_{\mathbb{G}, t-1}^{\text{OMDL}}[\mathcal{A}^*] = \text{forkAcc}(\mathcal{S}[\mathcal{A}])$$

(cf. steps 1-2 and (B.4)). Let $m = q_{ro} + 1$ be the number of distinct queries issued by \mathcal{A} to $\mathcal{O}^{\text{hash}}$ in the context of $\mathcal{S}[\mathcal{A}]$. The desired inequality follows from the Local Forking Lemma and Corollary 5.8. \square

Remark 5.19. (Known public shares) TODO

Algorithm 7 $\mathcal{A}^*(w_0, w_1, \dots, w_{t-1})$

- 1: Set $y^* \leftarrow w_0$
- 2: Set $x_i^* \leftarrow_{\S} \mathbb{Z}_q$, $1 \leq i \leq t-1$
- 3: $(b, \text{out}) \leftarrow \text{LocalFork}_{S[\mathcal{A}]}(y^*, x_1^*, \dots, x_{t-1}^*, w_1, \dots, w_{t-1})$
- 4: **if** $b = 0$
- 5: **return** \perp
- 6: Parse $(\tau, \tau^*) \leftarrow \text{out}$
- 7: Parse $(j, \{(u_i, c_i, s_i)\}_{i \in Q}) \leftarrow \tau$
- 8: Parse $(j, \{(u_i, c_i^*, s_i^*)\}_{i \in Q}) \leftarrow \tau^*$
- 9: $J \leftarrow Q \setminus \{j\}$
- 10: Compute the Lagrange interpolation coefficients $\{\lambda_i\}_{i \in Q}$
- 11: Compute

$$\mu_j \leftarrow \lambda_j \prod_{i \in J} c_i$$

- 12: **for** $i \in J$, compute

$$\mu_i \leftarrow \lambda_i \prod_{k \in Q \setminus \{i\}} c_k, \quad \mu_i^* \leftarrow \lambda_i \prod_{k \in Q \setminus \{i\}} c_k^*$$

- 13: $\{x_i\}_{i \in J} \leftarrow \text{Zip}(x_1^*, \dots, x_{t-1}^*; J)$
- 14: Query $r_i \leftarrow \mathcal{O}_{\mathbb{G}}^{\text{dlog}}(u_i)$, $i \in J$ // $t-1$ queries
- 15: Compute

$$x_j \leftarrow \frac{s_j - s_j^*}{c_j - c_j^*} + \frac{1}{\mu_j} \sum_{i \in J} \left(\mu_i (s_i - r_i - c_i x_i) - \mu_i^* (s_i^* - r_i - c_i x_i) \right)$$

- 16: Set $z_0 \leftarrow x_j$
 - 17: $(z_1, \dots, z_{t-1}) \leftarrow \text{Zip}^{-1}(\{r_i\}_{i \in J})$
 - 18: **return** z_0, z_1, \dots, z_{t-1}
-

APPENDIX A. SHAMIR'S SECRET SHARING

We list for reference some elementary facts about the sharing scheme introduced by Shamir [1]. Given a cyclic group $\mathbb{G} = \langle g \rangle$ of order q , we denote by $(\mathbb{Z}_q)_{t-1}[X]$ be the ring of polynomials of degree at most $t-1$ over \mathbb{Z}_q .

Definition A.1. The *Shamir (n, t) -sharing scheme*, $1 \leq t \leq n < q$, is the probabilistic algorithm $(x_1, \dots, x_n; y) \leftarrow D(x)$, $x \in \mathbb{Z}_q$ defined as

$$y \leftarrow g^x, \quad x_i \leftarrow f(x_i), \quad 1 \leq i \leq n, \quad f \leftarrow_{\$} (\mathbb{Z}_q)_{t-1}[X] \text{ with } f(0) = x$$

The space of possible Shamir (n, t) -sharings of x is denoted by

$$D_x = \{(x_1, \dots, x_n; y) \mid (x_1, \dots, x_n; y) \leftarrow D(x)\}$$

Remark A.2. (Reconstruction formula) Given $(x_1, \dots, x_n; y) \in D_x$, the condition

$$x = \sum_{i \in Q} \lambda_i x_i \iff |Q| \geq t \iff y = \prod_{i \in Q} y_i^{\lambda_i}, \quad (\text{A.3})$$

holds true, where $\{\lambda_i\}_{i \in Q}$ are the Lagrange interpolation coefficients and $y_i \equiv g^{x_i}$ are the public shares.

Remark A.4. ($t-1$ degrees of freedom) Given $J \subset \{1, \dots, n\}$ with $|J| = t-1$, pick $x_i^* \in \mathbb{Z}_q$, $i \in J$. Given $x \in \mathbb{Z}_q$, there is a unique $(x_1, \dots, x_n; y) \in D_x$ such that

$$x_i = x_i^*, \quad \forall i \in J$$

Remark A.5. (Security of Shamir's sharing) Given $x, x^* \in \mathbb{Z}_q$ and $J \subset \{1, \dots, n\}$ with $|J| = t-1$, the probability distributions

$$\{\{x_i\}_{i \in J} : (x_1, \dots, x_n; y) \leftarrow D(x)\}, \quad \{\{x_i\}_{i \in J} : (x_1, \dots, x_n; y) \leftarrow D(x^*)\}$$

are indistinguishable.

APPENDIX B. LOCAL FORKING LEMMA

The Forking Lemma was introduced by Pointcheval and Stern ([4], [5]) as part of the rewinding methodology for proving the security of Fiat-Shamir based signature schemes in the random oracle; it was further elaborated as General Forking Lemma by Bellare and Neven [12], who abstracted away the signature specific details. The Local Forking Lemma is a variance by [11] which dispenses with rewinding completely, allowing to fork a (deterministic) process by reprogramming the oracle at a single location as opposed to every location after the fork.

Consider a deterministic algorithm of the form

$$\mathcal{S} : \Theta \times C^m \rightarrow \{0, 1, \dots, m\} \times T, \quad (k, \tau) \leftarrow \mathcal{S}(\theta; h_1, \dots, h_m), \quad (\text{B.1})$$

running internally a hash oracle with preallocated hashes h_1, \dots, h_m in respective order (cf. Section 5.3). Think of θ as the random coins, k as the *main* output and τ as the *side* output. If $k \geq 1$, the main output may be thought of as indicating some peculiarity of the k -th distinct oracle query regarding the side output, while the case $k = 0$ is interpreted as failure. We denote by

$$\text{acc}(\mathcal{S}) = \Pr[k \geq 1 : \left\{ \begin{array}{l} (k, \tau) \leftarrow \mathcal{S}(\theta; h_1, \dots, h_m), \quad \theta \leftarrow_{\$} \Theta \\ h_\nu \leftarrow_{\$} C, \quad 1 \leq \nu \leq m \end{array} \right\}] \quad (\text{B.2})$$

the probability that \mathcal{S} accepts taken over the possible coins and oracle outputs.

Definition B.3. The *local fork* of \mathcal{S} is Algorithm 8.

Algorithm 8 LocalFork $\mathcal{S}(\theta)$

```
1:  $h_\nu \leftarrow_{\$} C, \ 1 \leq \nu \leq m$ 
2:  $(k, \tau) \leftarrow \mathcal{S}(\theta; h_1, \dots, h_m)$ 
3: if  $k = 0$ 
4:   return  $(0, \perp)$ 
5:  $h_k^* \leftarrow_{\$} C$ 
6:  $(k^*, \tau^*) \leftarrow \mathcal{S}(\theta; h_1, \dots, h_{k-1}, h_k^*, h_{k+1}, \dots, h_m)$ 
7: if  $k^* = k \wedge h_k \neq h_k^*$ 
8:   return  $(1, (\tau, \tau^*))$ 
9: return  $(0, \perp)$ 
```

That is, if \mathcal{S} accepts, the fork reruns it with the same input except for the k -th hash, which is afresh uniformly sampled; otherwise it aborts. When the second execution completes, the fork checks if \mathcal{S} accepts again under the constraint $h_k^* \neq h_k$ and returns both side outputs; otherwise, it aborts. We denote by

$$\text{forkAcc}(\mathcal{S}) = \Pr[b = 1 : (b, \text{out}) \leftarrow \text{LocalFork}_{\mathcal{S}}(\theta), \theta \leftarrow_{\$} \Theta] \quad (\text{B.4})$$

the probability that the local fork does not abort taken over the possible random coins of \mathcal{S} . The local forking lemma asserts the following lower bound.

Lemma B.5. (*Local Forking Lemma*, [11]) *For every deterministic algorithm \mathcal{S} as in (B.1) issuing m distinct queries to a hash oracle with tape $[h_1, \dots, h_m]$,*

$$\text{forkAcc}(\mathcal{S}) \geq \frac{1}{m} \text{acc}(\mathcal{S})^2 \quad (\text{B.6})$$

REFERENCES

- [1] Adi Shamir. How to Share a Secret. Commun. ACM 1979.
- [2] Claus-Peter Schnorr. Efficient Signature Generation by Smart Cards. Journal of Cryptology 1991, p. 161-174
- [3] C. Komlo and I. Goldberg. FROST: Flexible Round-Optimized Schnorr Threshold Signatures. SAC 2020.
- [4] David Pointcheval, Jacques Stern. Security proofs for Signature Schemes. Advances in Cryptology - EUROCRYPT '96, International Conference on the Theory and Application of Cryptographic Techniques, Saragossa, Spain, May 12-16, 1996, Proceeding.
- [5] David Pointcheval, Jacques Stern. Security Arguments for Digital Signatures and Blind Signatures. Journal of Cryptology, 2000, p. 361-396
- [6] Douglas R. Stinson, Reto Strohle. Provably Secure Distributed Schnorr Signatures and a (t, n) Threshold Scheme for Implicit Certificates. ACISP 2001.
- [7] Marcel Keller, Gert Læssøe Mikkelsen, Andy Rupp: Efficient Threshold Zero-Knowledge with Applications to User-Centric Protocols. ICITS 2012: 147-166
- [8] Maurer, U.M.: Unifying zero-knowledge proofs of knowledge. In: Preneel, B. (ed.) AFRICACRYPT. Lecture Notes in Computer Science, vol. 5580, pp. 272–286. Springer (2009)
- [9] Mihir Bellare, Adriana Palacio. GQ and Schnorr Identification Schemes: Proofs of Security against Impersonation under Active and Concurrent Attacks. CRYPTO 2002.
- [10] Mihir Bellare, Chanathip Namprempre, David Pointcheval, and Michael Semanko. The one-more-RSA-inversion problems and the security of Chaum’s blind signature scheme. Journal of Cryptology, 16(3):185–215, June 2003.

- [11] Mihir Bellare, Wei Dai, and Lucy Li. The local forking lemma and its application to deterministic encryption. In S. D. Galbraith and S. Moriai, editors, *Advances in Cryptology - ASIACRYPT 2019 - 25th International Conference on the Theory and Application of Cryptology and Information Security*, Kobe, Japan, December 8-12, 2019, *Proceedings, Part III*, volume 11923 of *Lecture Notes in Computer Science*, pages 607–636. Springer, 2019.
- [12] Mihir Bellare and G. Neven. Multi-signatures in the plain public-key model and a general forking lemma. In A. Juels, R. N. Wright, and S. D. C. di Vimercati, editors, *Proceedings of the 13th ACM Conference on Computer and Communications Security, CCS 2006*, Alexandria, VA, USA, October 30 - November 3, 2006, pages 390–399. ACM, 2006.